

Software tools for calculating materials properties in high-throughput (pymatgen, atomate, FireWorks)

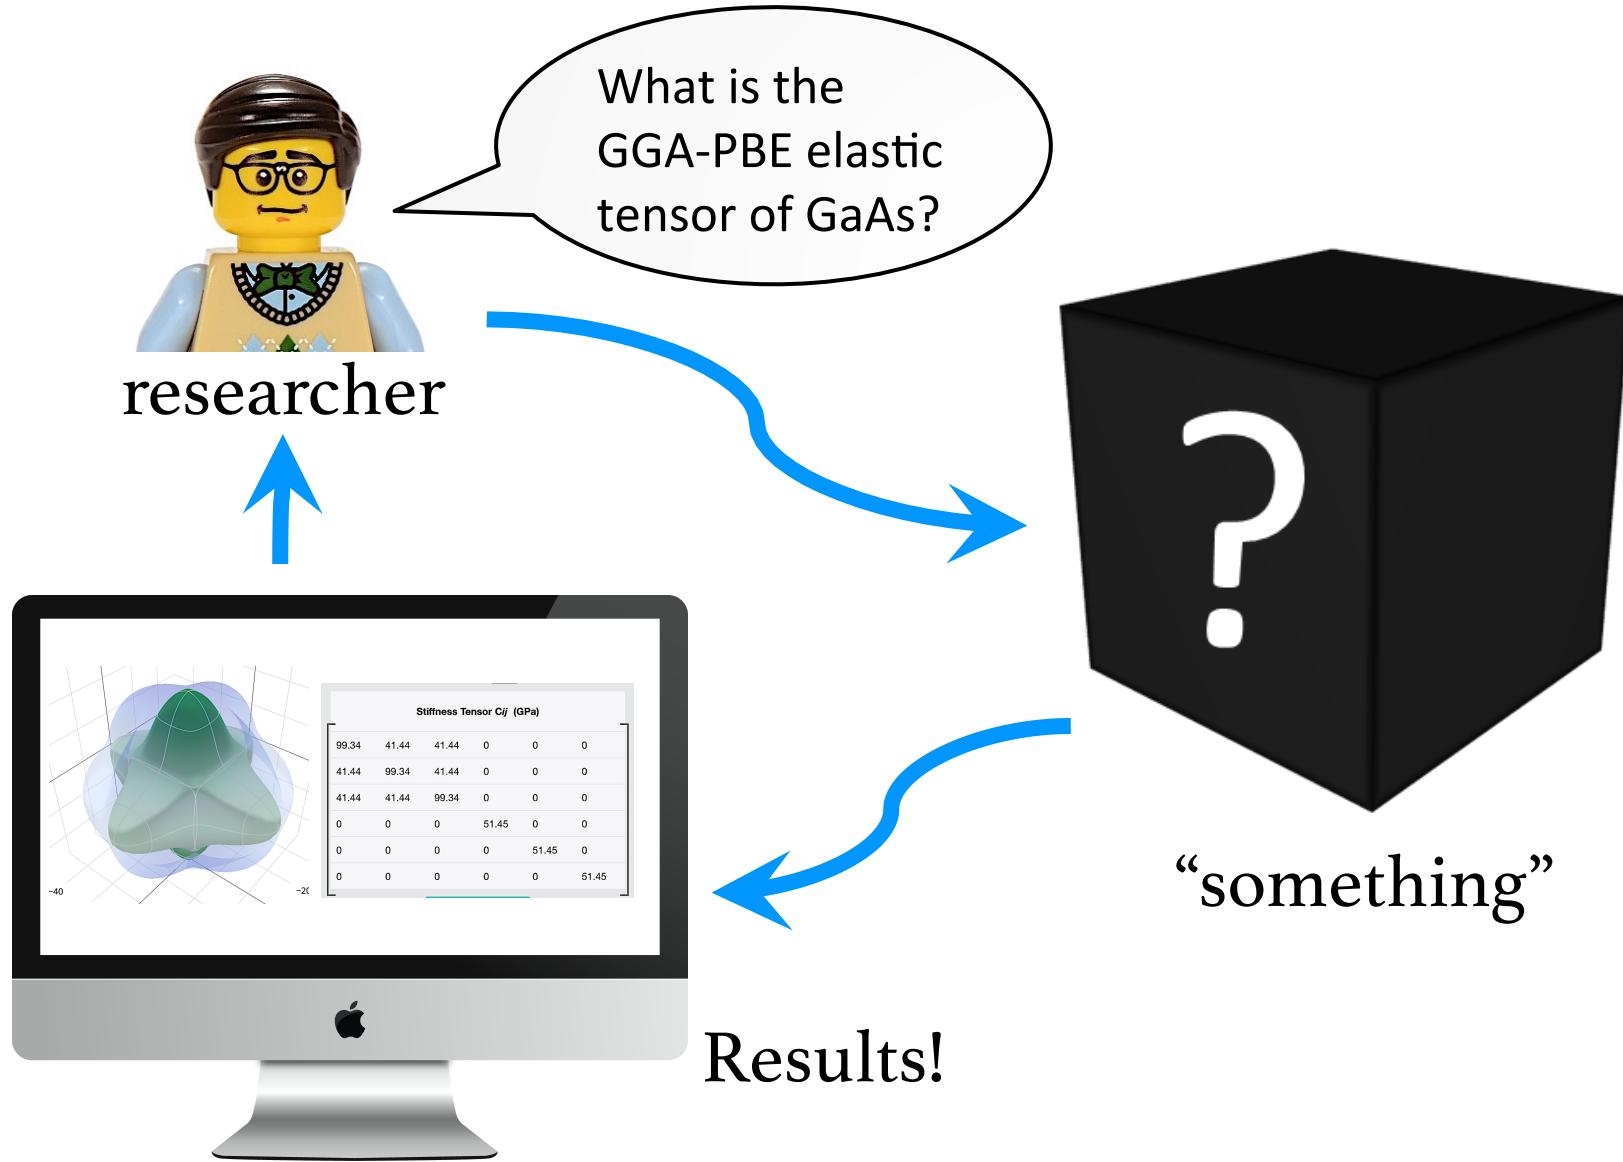
Anubhav Jain
Energy Technologies Area
Lawrence Berkeley National Lab
Berkeley, CA

MP workshop 2017

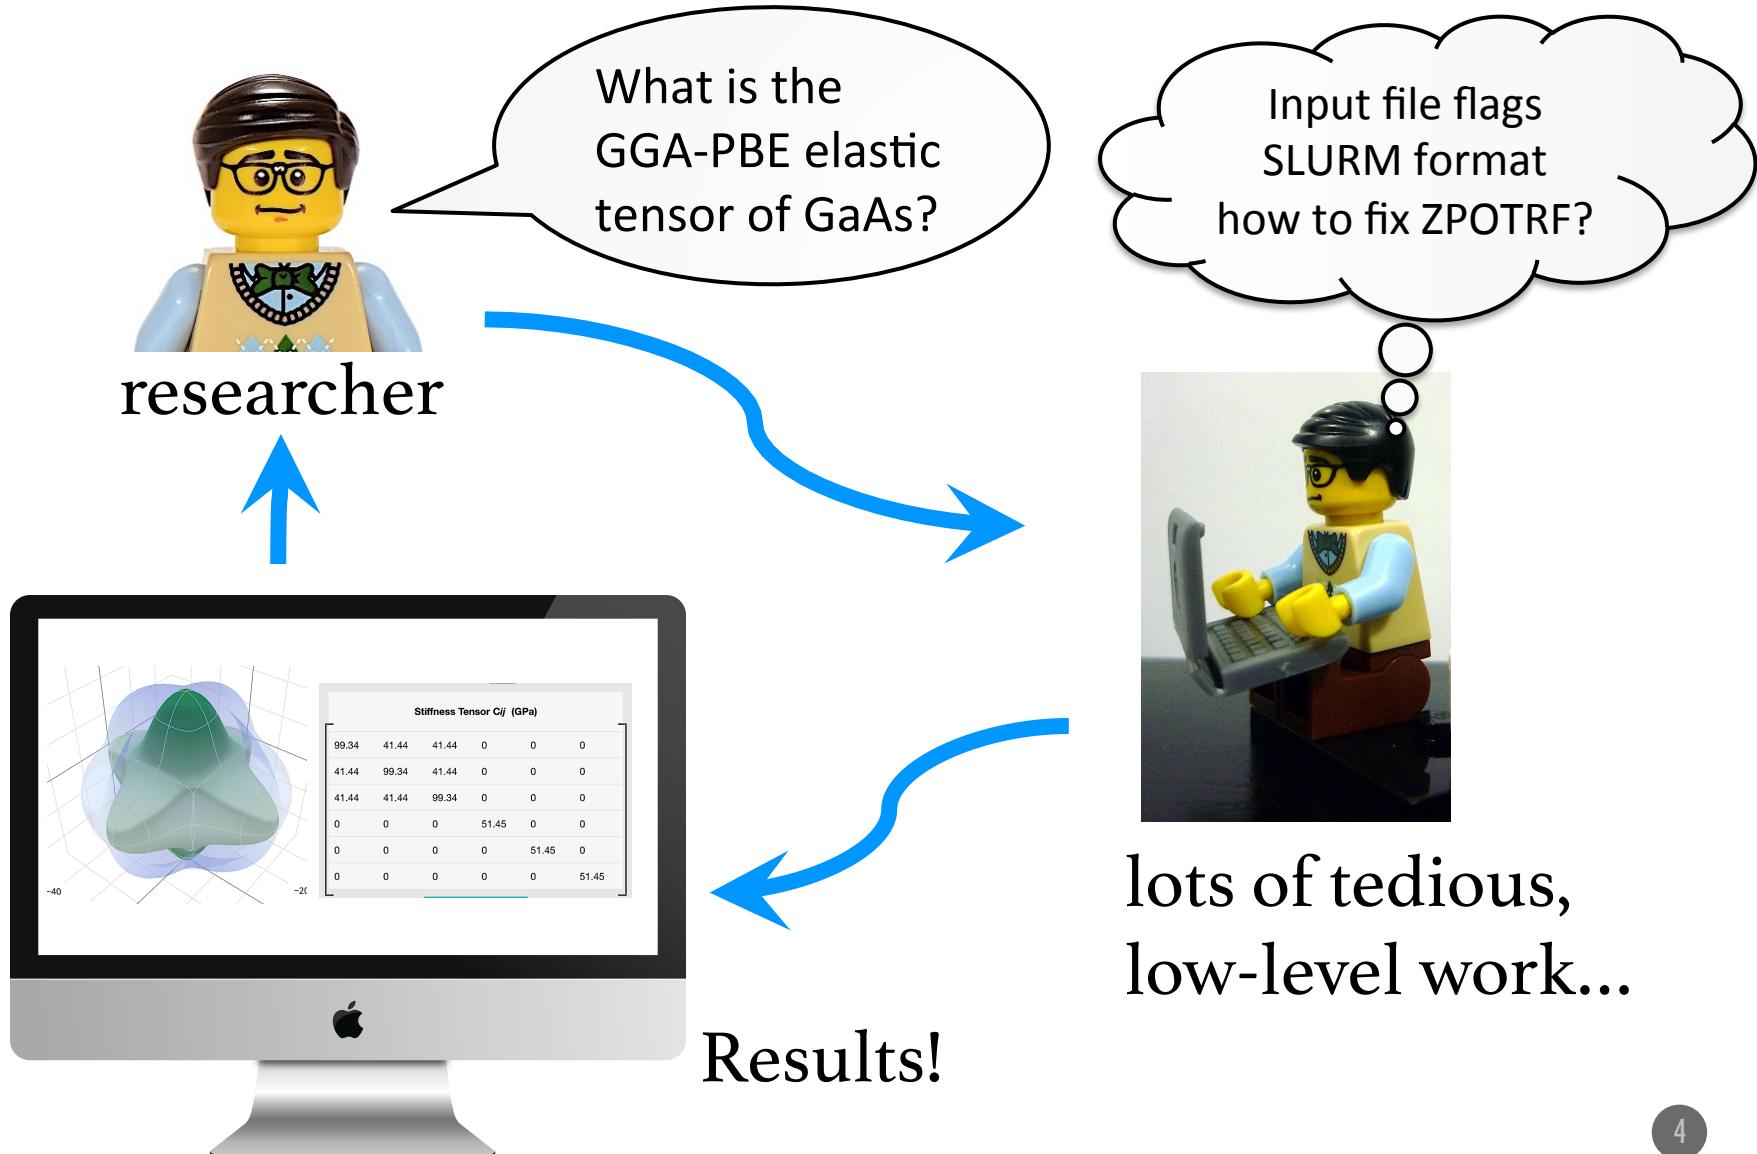
“Civilization advances by extending the number of important operations which we can perform without thinking about them.”

- Alfred North Whitehead

A “black-box” view of performing a calculation



Unfortunately, the inside of the “black box” is usually tedious and “low-level”



Calculations are labor intensive!

- Some of the things to do include:
 - set up the structure coordinates
 - write input files, double-check all the flags
 - copy to supercomputer
 - submit job to queue
 - deal with supercomputer headaches
 - monitor job
 - fix error jobs, resubmit to queue, wait again
 - repeat process for subsequent calculations in workflow
 - parse output files to obtain results
 - copy and organize results, e.g., into Excel
- This is often repeated for each and every run!



All the low-level steps can lead to errors!

Let's take a look at two alternate universes:

1



you



have coffee



copy files from
previous simulation



edit 5 lines



run simulation,
analyze data

2



you



forget coffee



copy files from
previous simulation



edit 4 lines
but forget
LHFCALC=F



run simulation,
looks fine at first,
in a month you
discover it was wrong

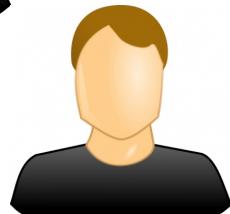
It is too easy to end up in
the wrong timeline!

Today, it is difficult to learn and apply several computational procedures due to steep learning curve

Because of the multiple low-level steps that all need to be done correctly to apply computational methods, there is often a single group “expert” for each technique



“Alice knows how to do charged defect calculations.”

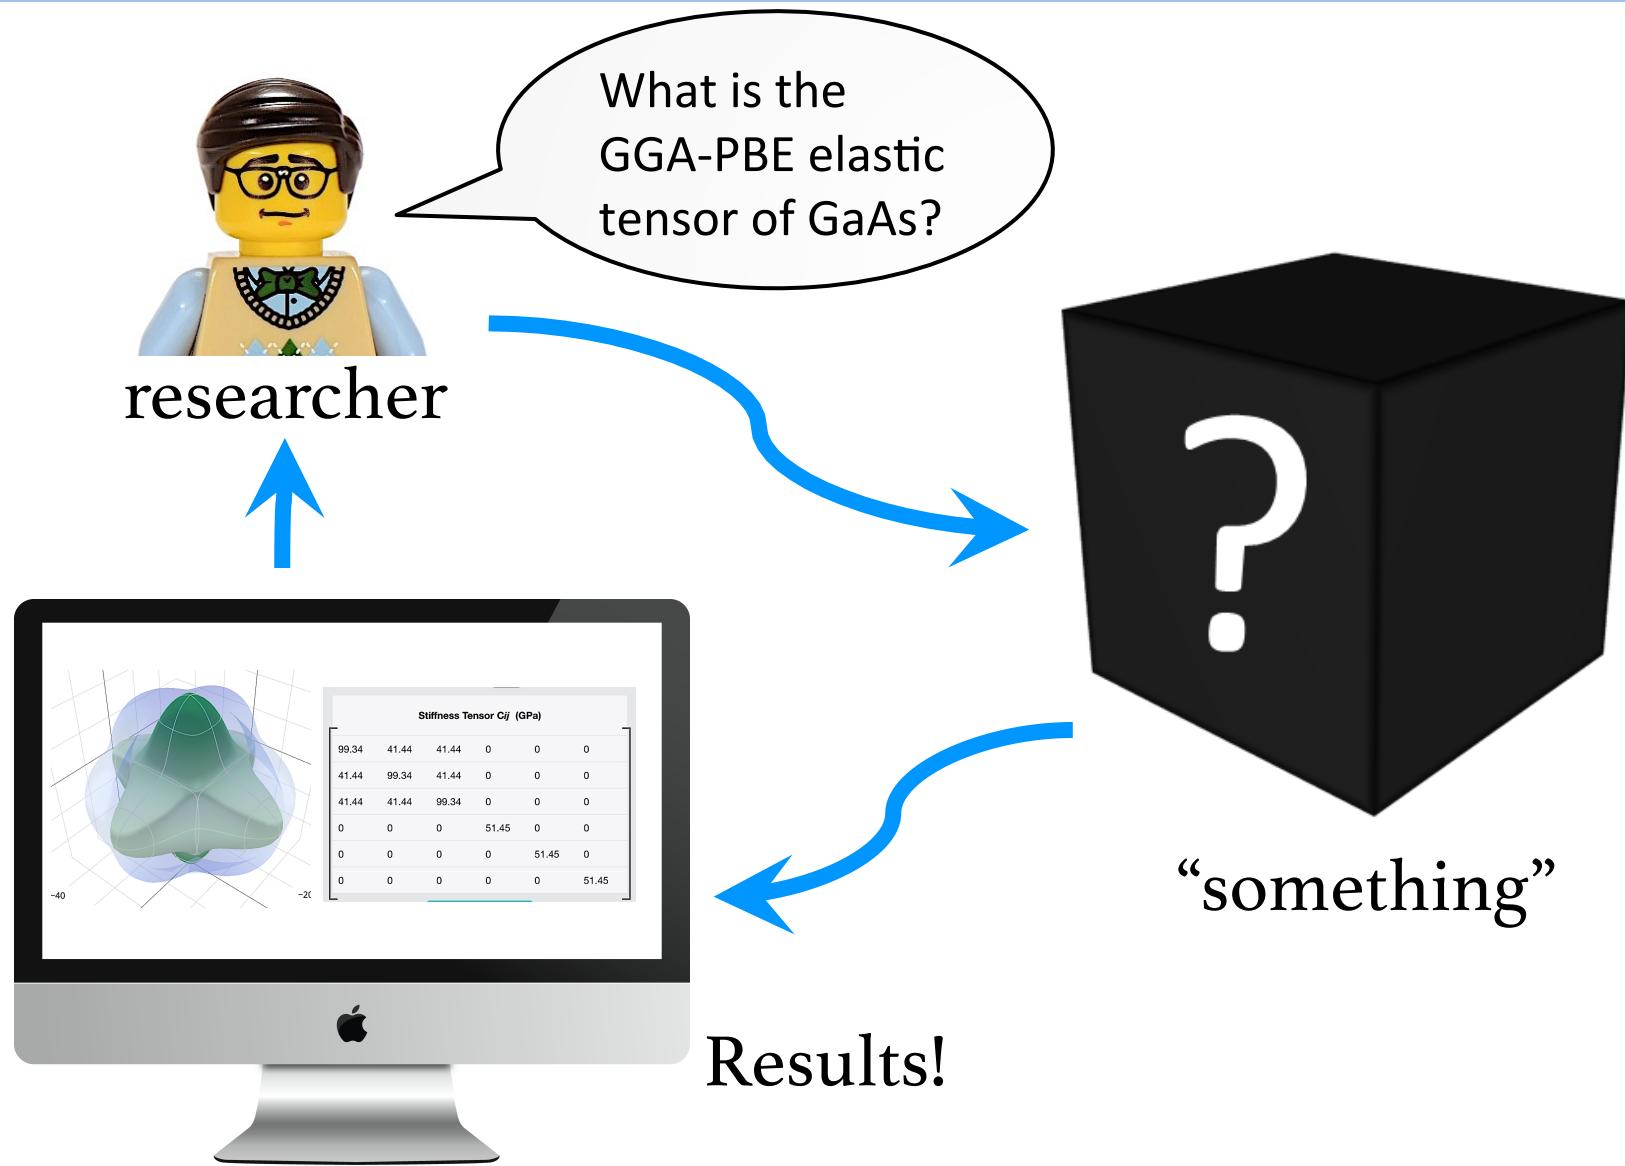


“Bob is the one who can properly converge GW runs.”

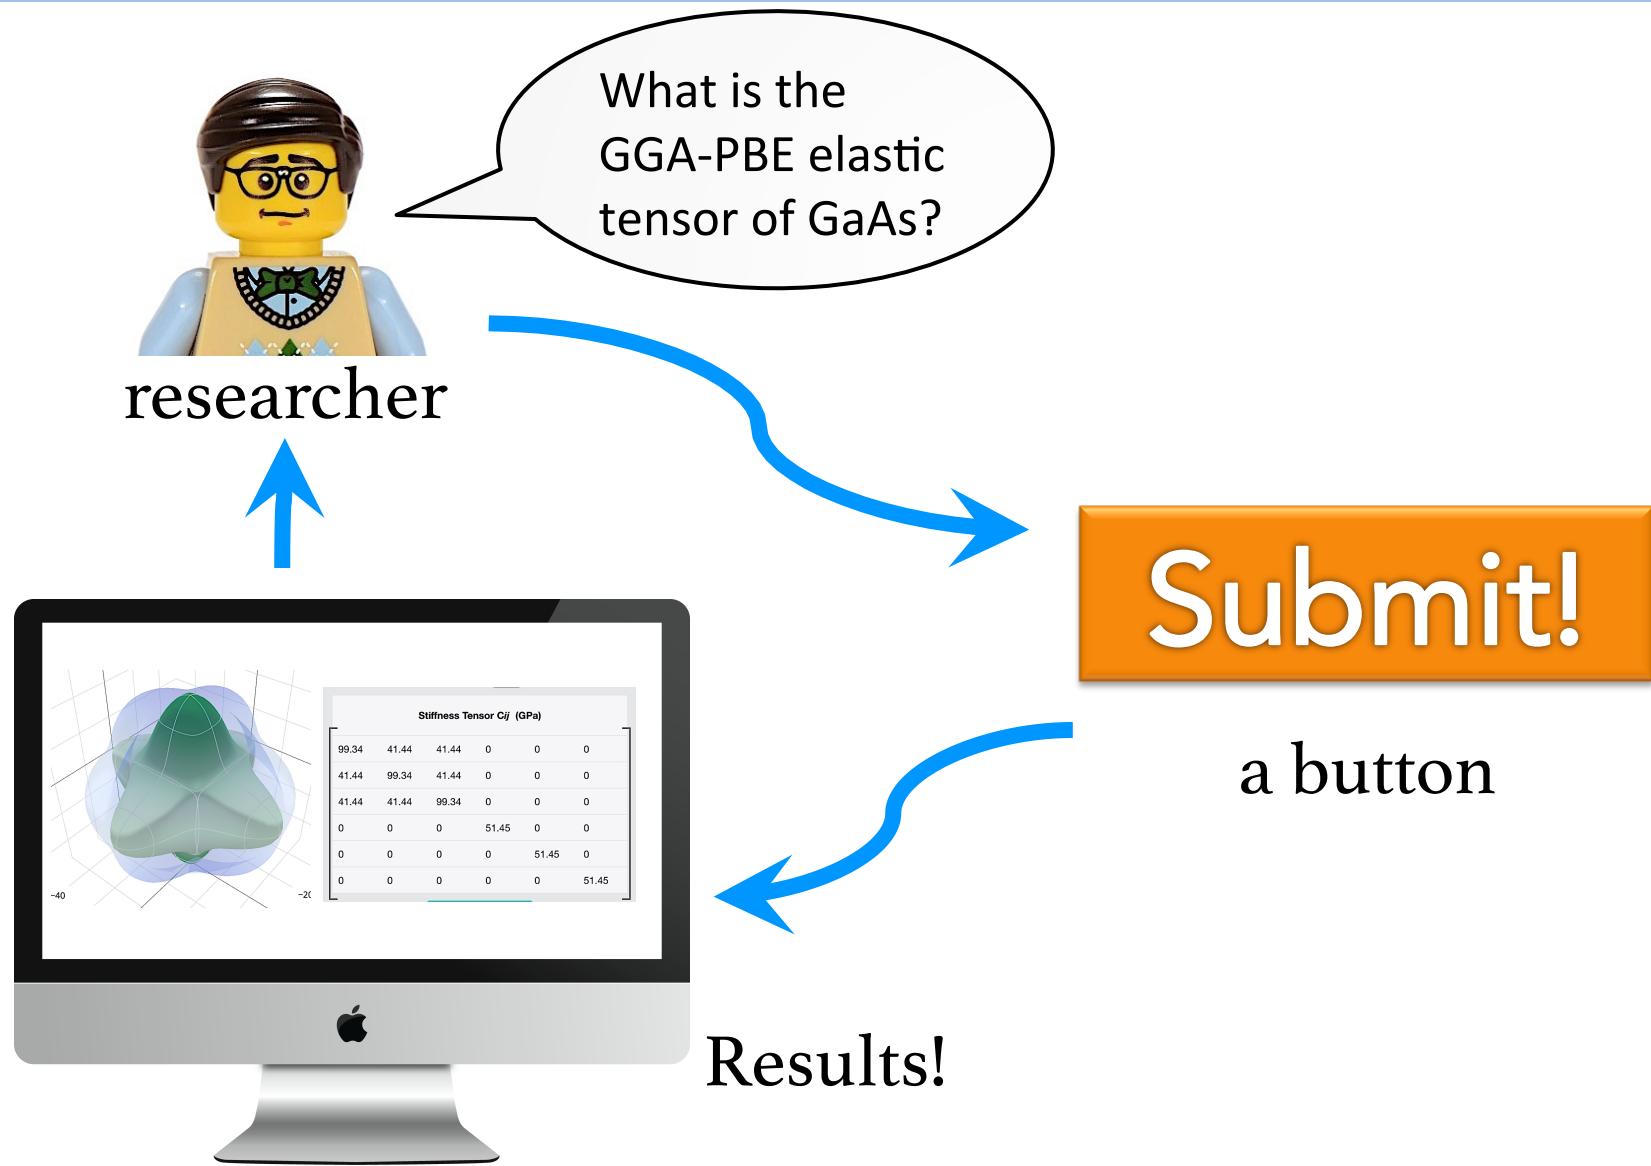


“Olga has all the scripts for phonon calculations.”

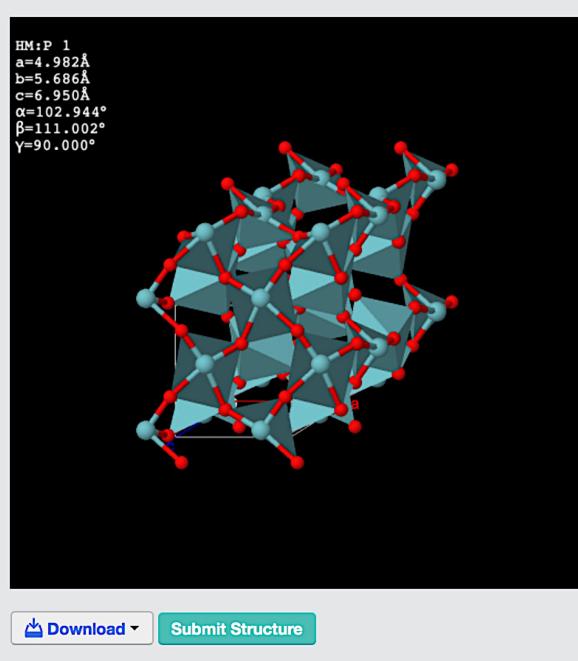
What would be a better way?



What would be a better way?



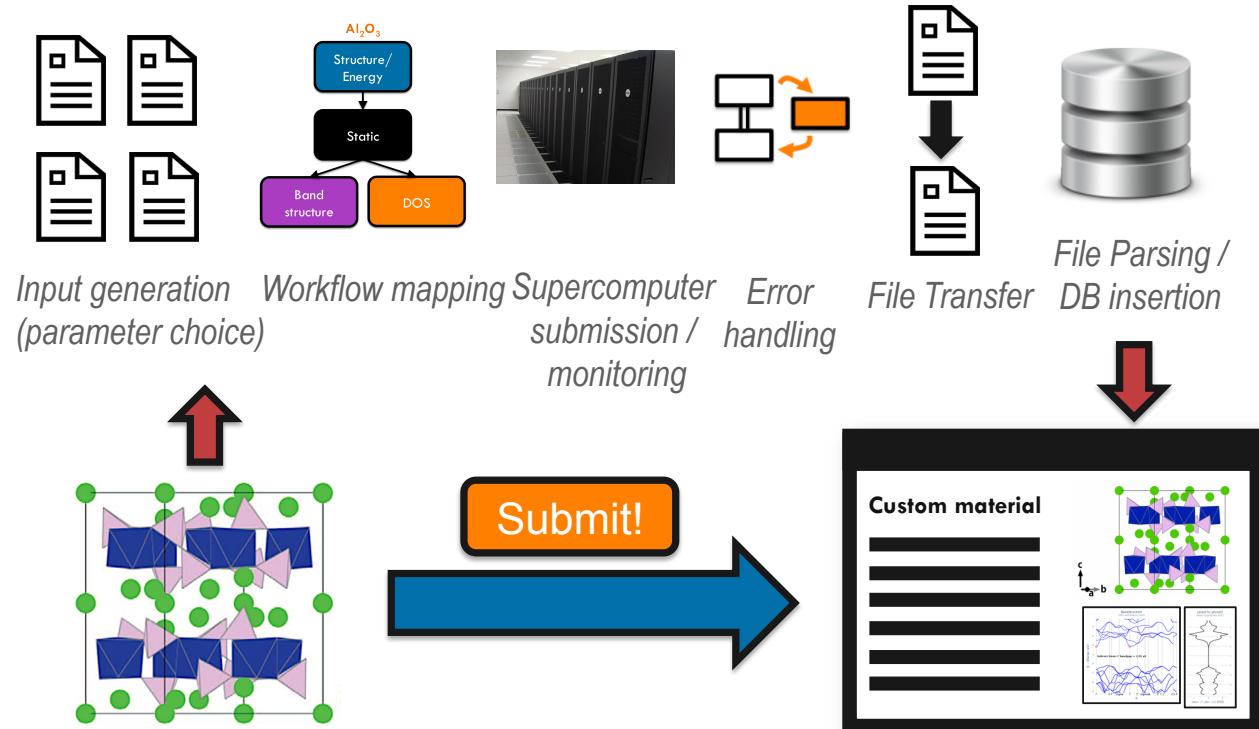
MPComplete on Materials Project works this way



www.materialsproject.org

“Crystal Toolkit”

Anyone can find, edit,
and submit (suggest)
structures



Currently, this feature is available for:

- structure optimization
- band structures
- elastic tensors

What if you want to do something more complicated?



- Start with all the known binary oxides
- Substitute oxygen with sulfur
- Re-optimize the structure and compute the band structure
- Plot the band gap of oxide vs the sulfide
- Rerun compounds with heavy elements with spin-orbit coupling enabled

Now, a simple button won't do it anymore.

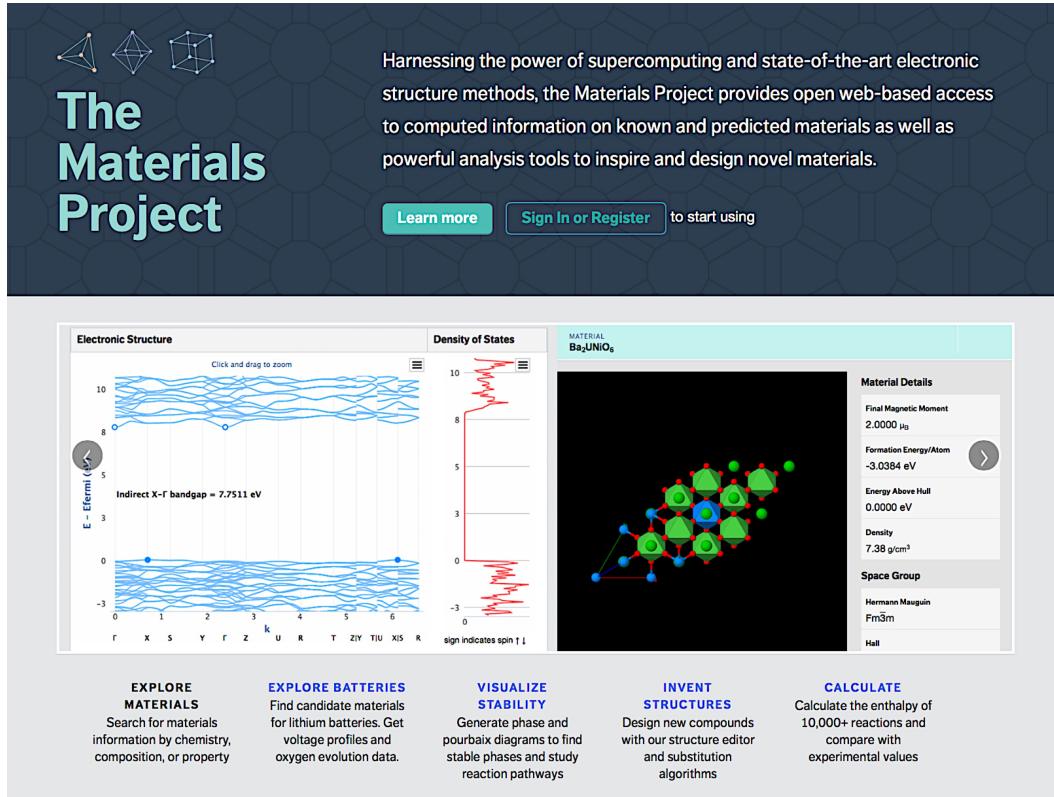
What you want is a high-level language for defining, running, and executing simulations.

With such a capability, one could do many applications.

Example: The Materials Project database

The Materials Project (<http://www.materialsproject.org>)

free and open



Jain*, Ong*, Hautier, Chen, Richards, Dacek, Cholia, Gunter, Skinner, Ceder, and Persson, APL Mater., 2013, 1, 011002. *equal contributions

>30,000 registered users around the world

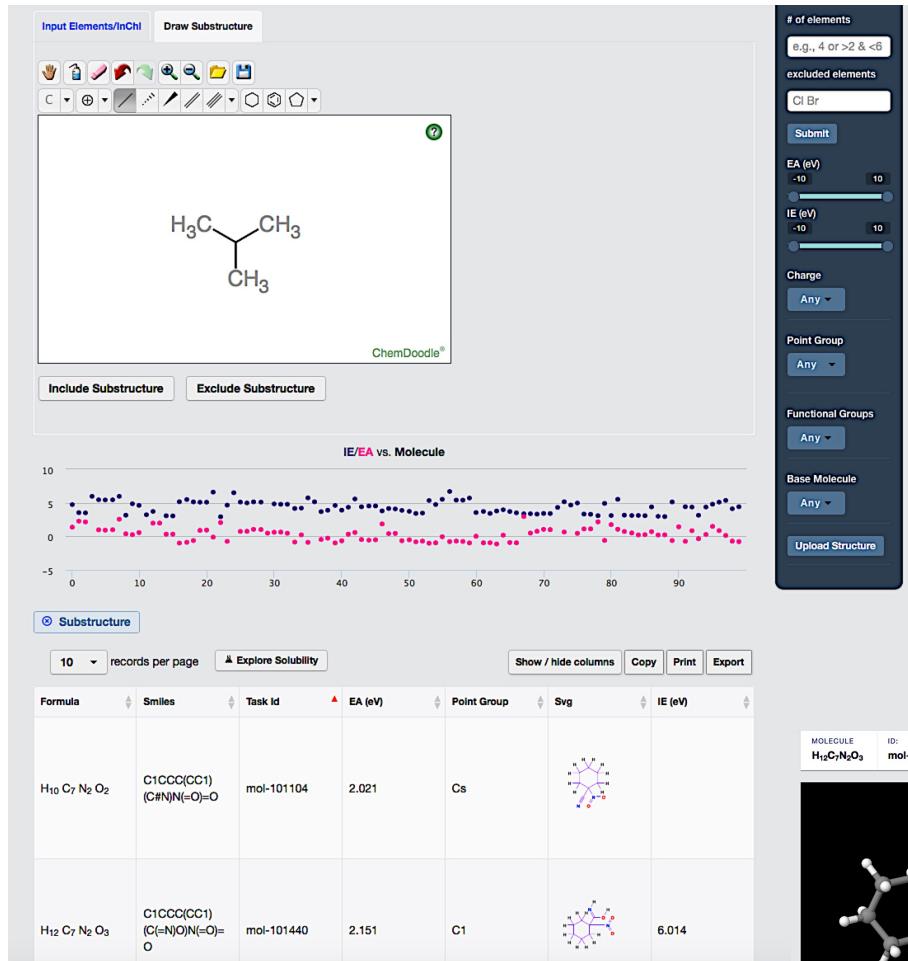
>65,000 compounds calculated

Data includes

- thermodynamic props.
- electronic band structure
- aqueous stability (E-pH)
- elasticity tensors
- piezoelectric tensors

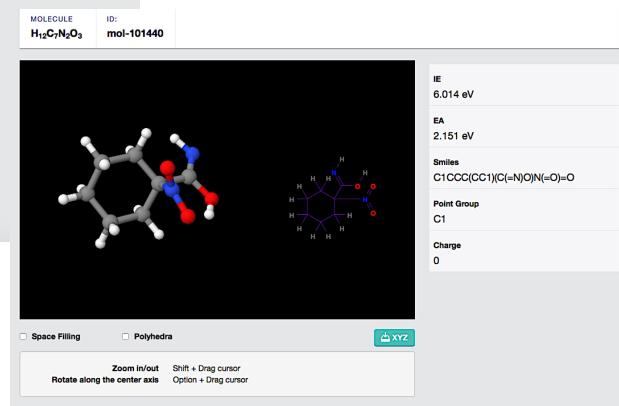
>75 million CPU-hours invested = massive scale!

Example: The Electrolyte Genome



data on ~22,000 molecules
(mainly geometry + IP/EA via full adiabatic calcs)

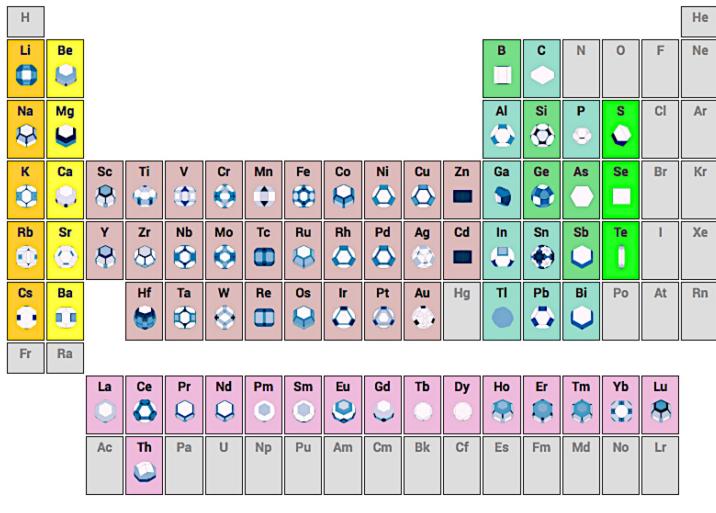
Also deployed on the Materials Project web site



L. Cheng, R.S. Assary, X. Qu, A. Jain, S.P. Ong, N.N. Rajput, et al., J. Phys. Chem. Lett. 6 (2015) 283–291.

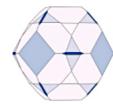
X. Qu, A. Jain, N.N. Rajput, L. Cheng, Y. Zhang, S.P. Ong, et al., Comput. Mater. Sci. 103 (2015) 56–67.

Example: Crystalium (Ong / Persson)



Ag polymorph #0 (mp-124, Fm $\bar{3}m$, $E_{full}=0.000$ eV/atom)

Cycle through different views
(111) (200) (210) (300) (310)
(001) (010) (100) (020) (112)



Miller Indices	γ		Area fraction
	J/m ²	eV/ \AA^2	
(111)	0.77	0.048	0.04
(322)	0.77	0.048	0.65
(332)	0.79	0.049	0.10
(100)	0.81	0.051	0.27
(221)	0.82	0.051	0.00
(331)	0.85	0.053	0.00
(321)	0.86	0.054	0.00
(311)	0.86	0.054	0.00
(110) [*]	0.87	0.054	0.02
(211)	0.87	0.054	0.00
(110)	0.87	0.055	0.02
(320)	0.89	0.055	0.00
(310)	0.89	0.056	0.00
(210)	0.90	0.056	0.00

* indicates reconstructed surface.

Help and definitions

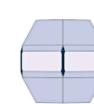
Weighted surf. energy $\bar{\gamma}$ = 0.79 J/m² (0.049 eV/ \AA^2)

Shape factor η = 5.06

Surf. energy anisotropy δ_γ = 0.023

Ag polymorph #1 (mp-989737, R $\bar{3}m$, $E_{full}=0.003$ eV/atom)

Cycle through different views
(001) (010) (100) (020) (112)
(011) (012) (111) (110) (120)



Miller Indices	γ		Area fraction
	J/m ²	eV/ \AA^2	
(011)	0.74	0.046	1.00
(112)	0.85	0.053	0.00
(121)	0.87	0.054	0.00

* indicates reconstructed surface.

Help and definitions

Weighted surf. energy $\bar{\gamma}$ = 0.74 J/m² (0.046 eV/ \AA^2)

Shape factor η = 8.19

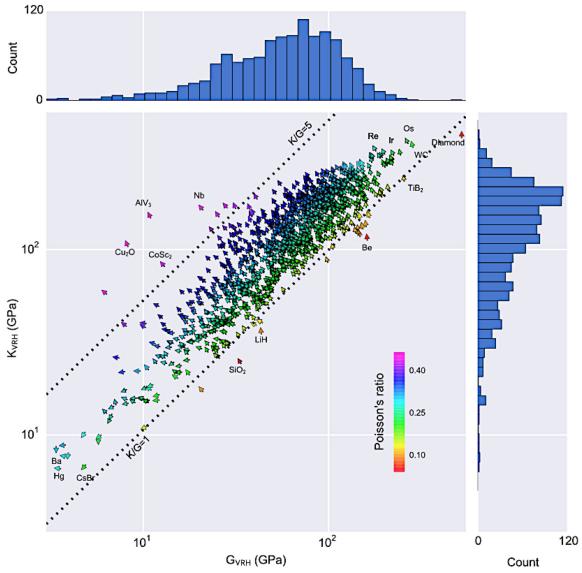
Surf. energy anisotropy δ_γ = 0.002

<http://crystalium.materialsvirtuallab.org>

surface energies for 142 polymorphs of 72 elements + rotatable Wulff shapes

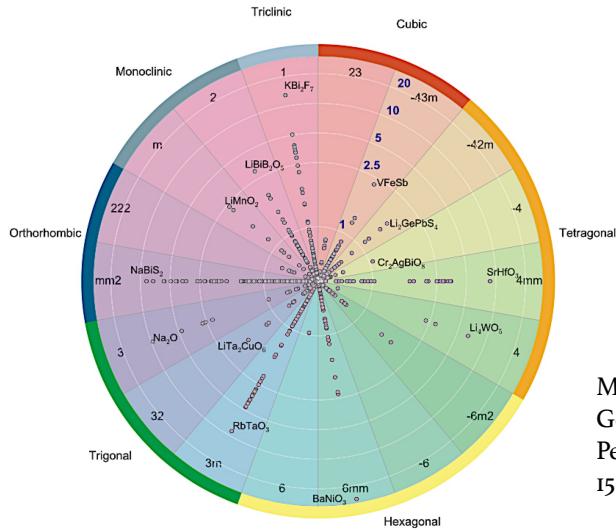
computed & maintained by the Ong group (UC San Diego) with support from Persson Group (UC Berkeley)

Example: Rapid data generation



>4500 elastic tensors

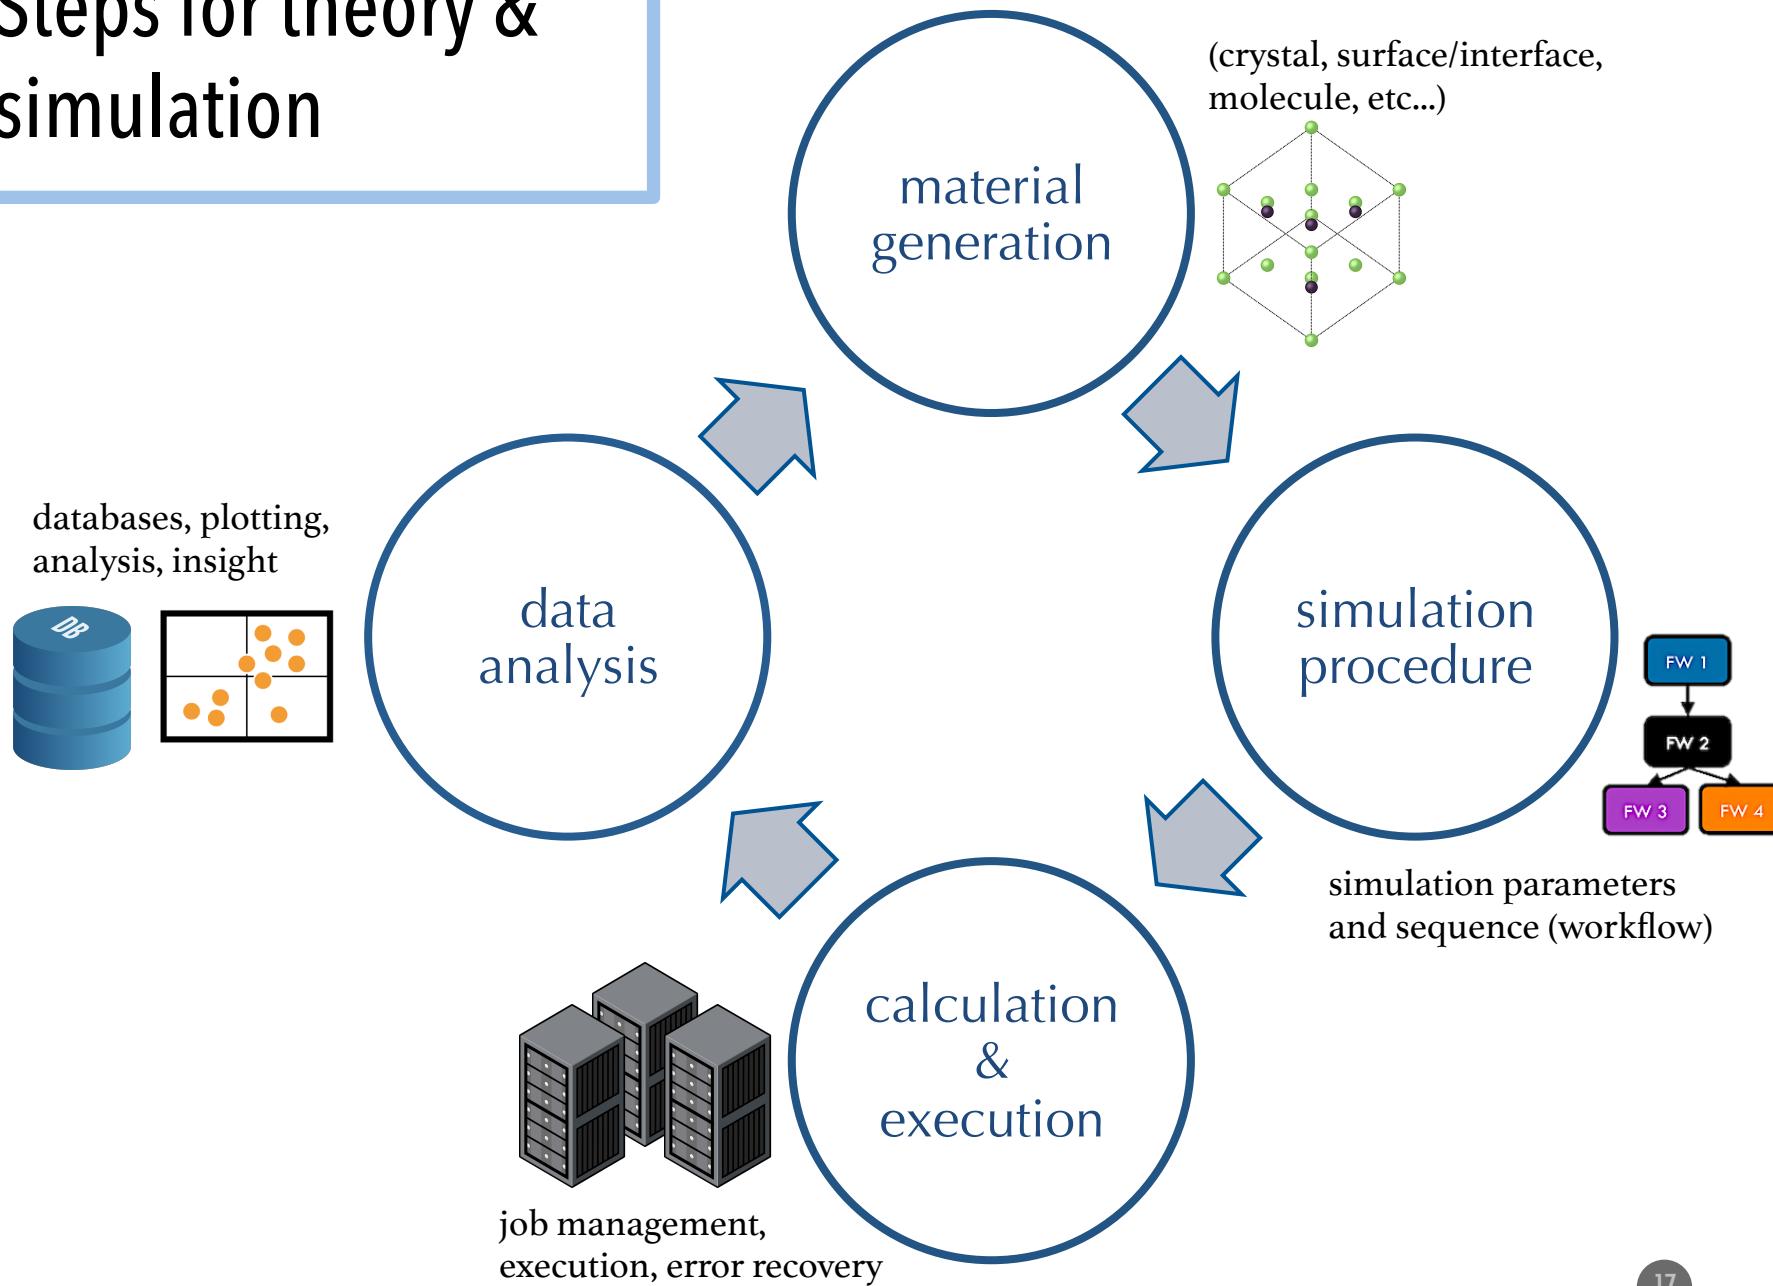
M. De Jong, W. Chen, T. Angsten, A. Jain, R. Notestine, A. Gamst, M. Sluiter, C. K. Ande, S. Van Der Zwaag, J. J. Plata, C. Toher, S. Curtarolo, G. Ceder, K. a Persson, and M. Asta, Sci. Data, 2015, 2, 150009.



How can we design a simple yet powerful tool for performing simulations?

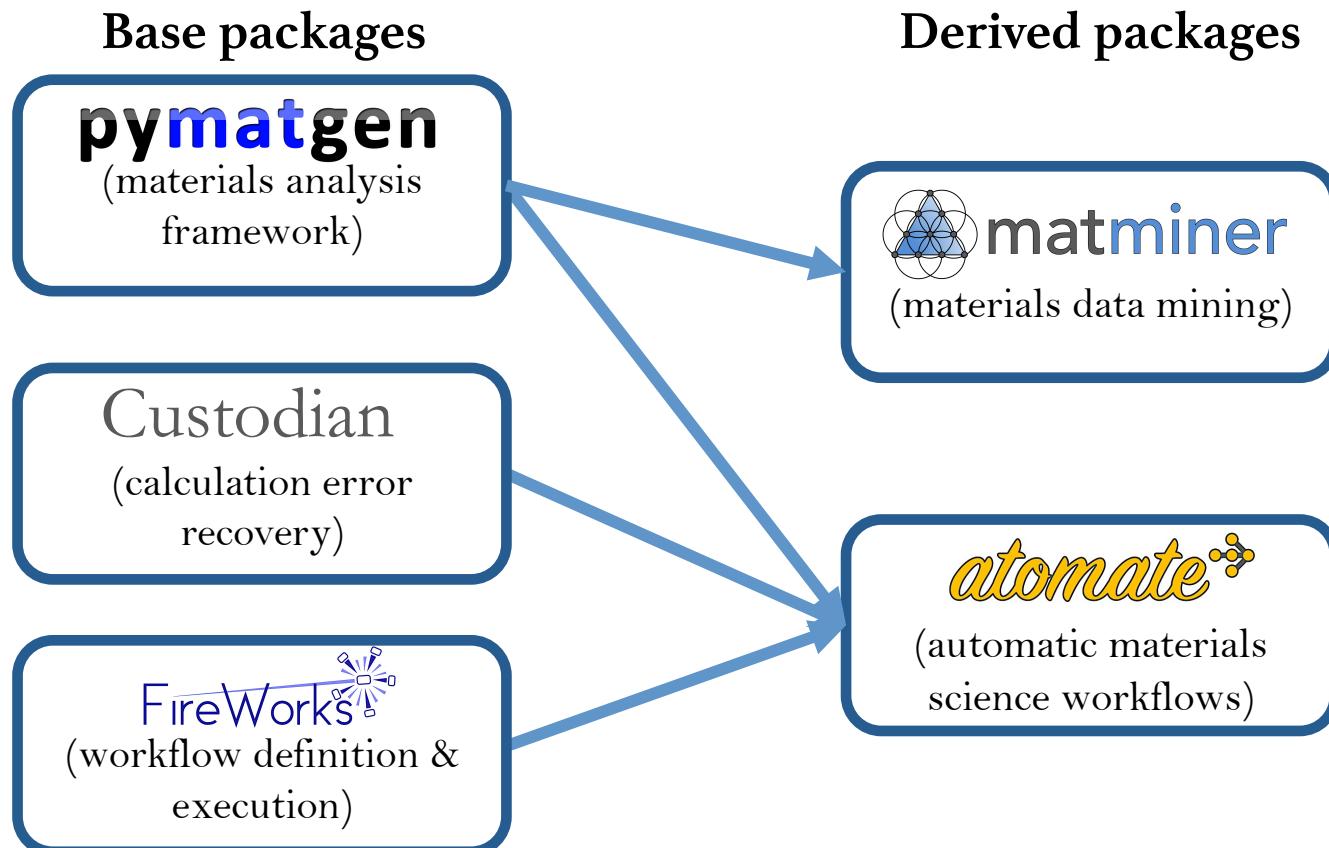
- First, we'll go through the four high-level steps in doing a theory study
- In each step, we'll talk about how software tools can help make you more productive and help you concentrate on things that require your brainpower, not mindless labor (e.g., manually resubmitting failed jobs)

Steps for theory & simulation

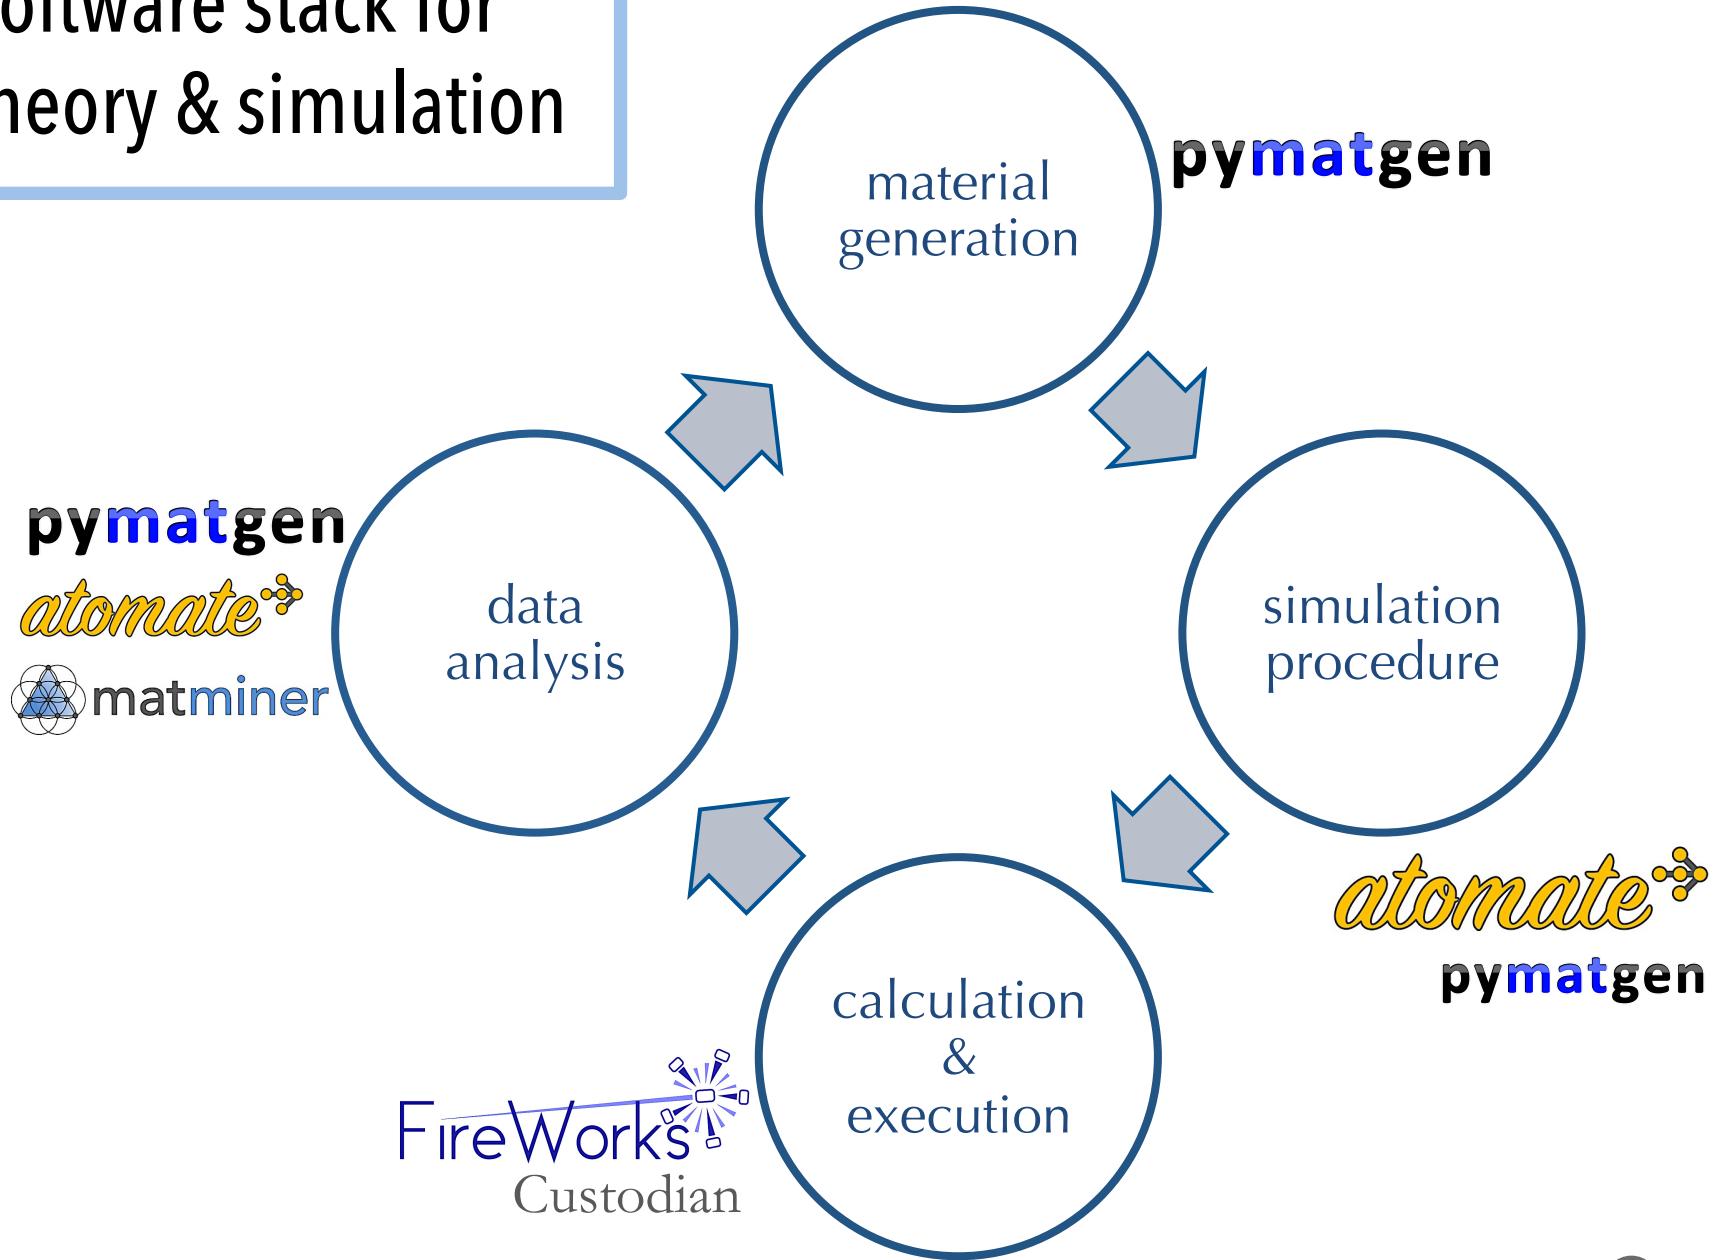


Software technologies for theory / calculation

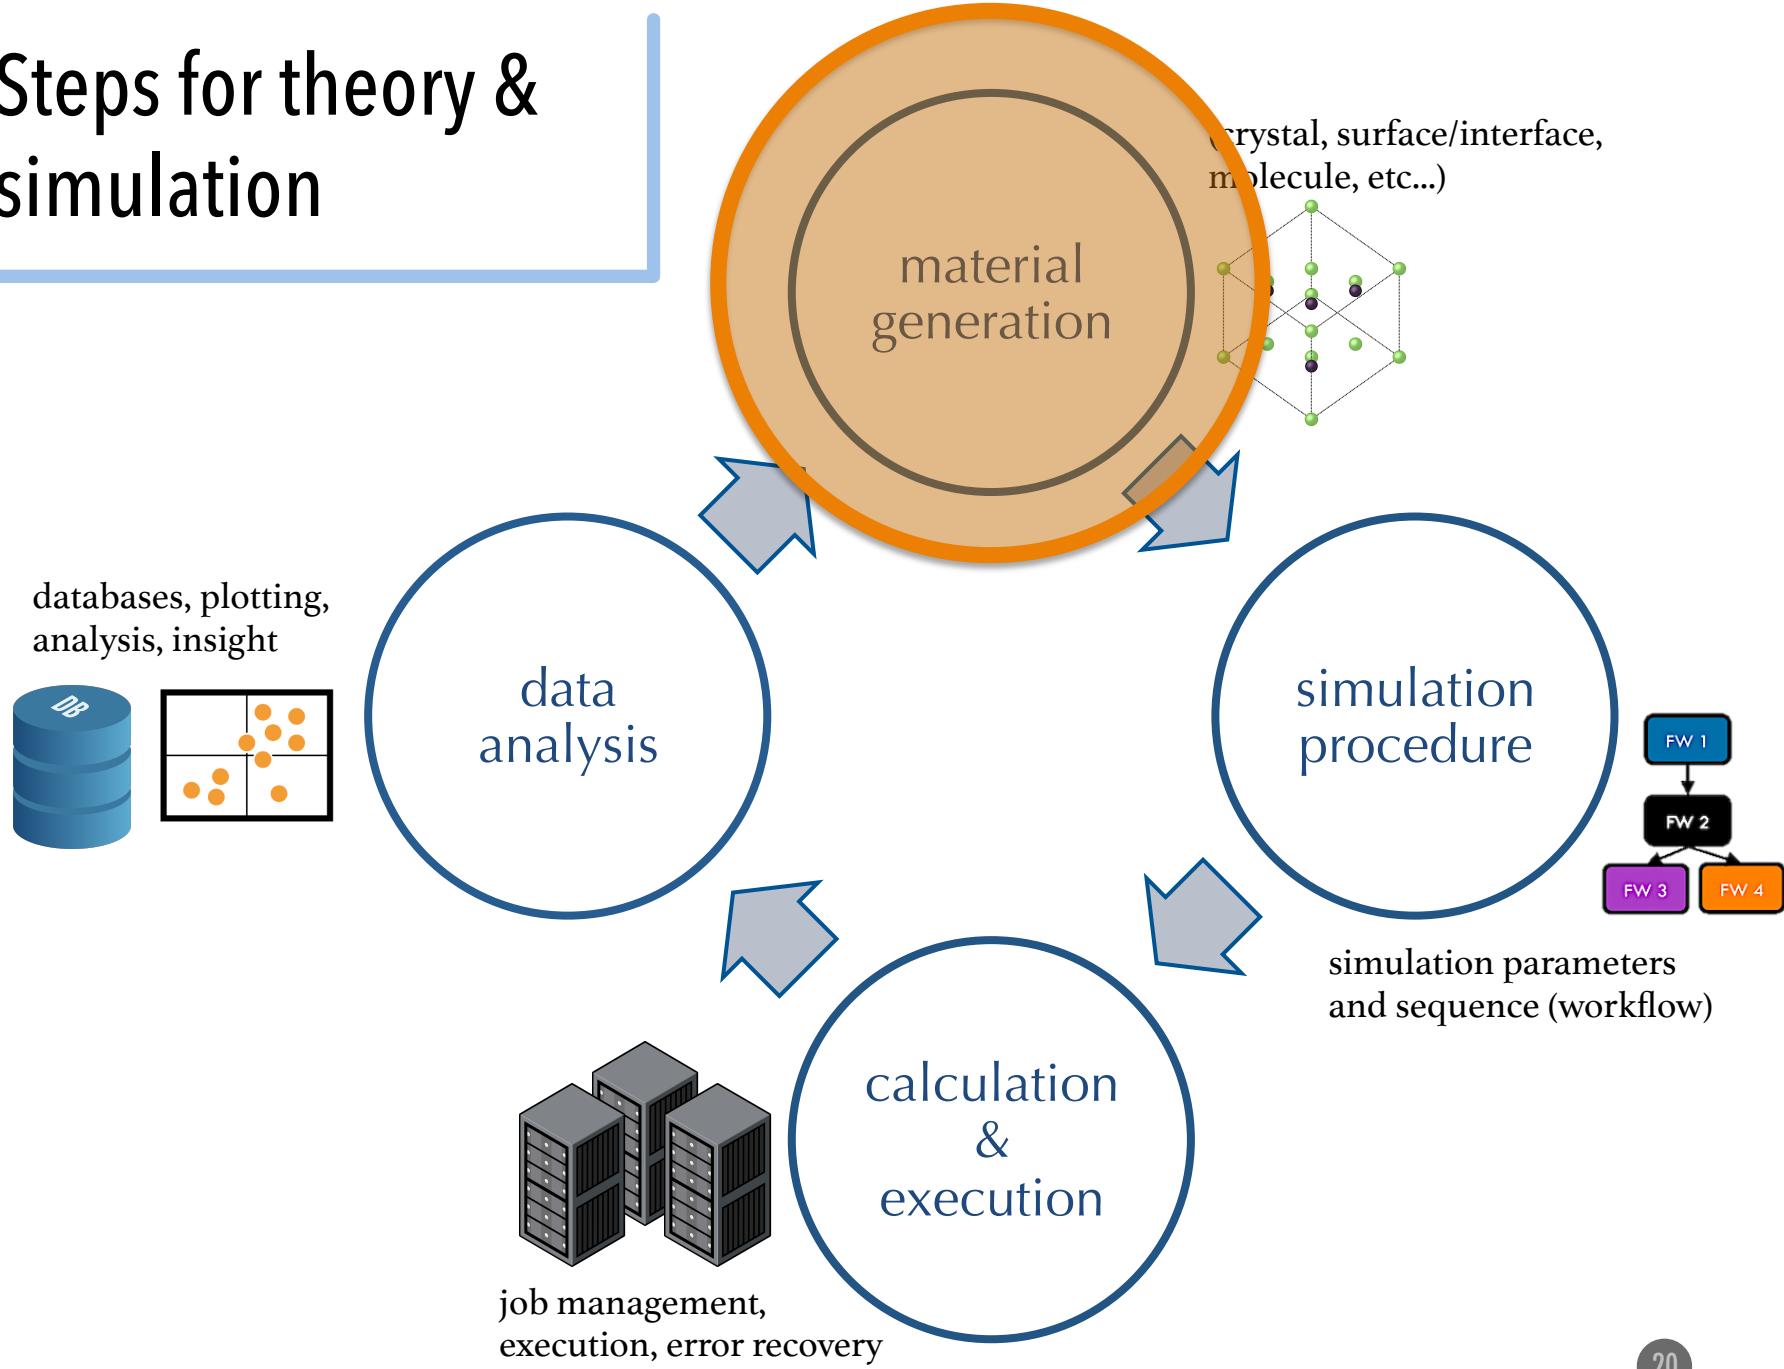
These are all open-source:



Software stack for theory & simulation

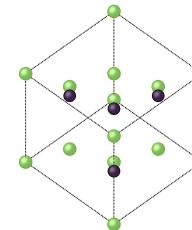


Steps for theory & simulation



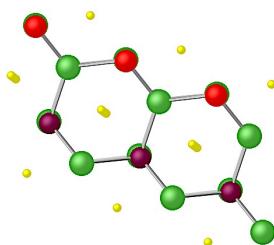
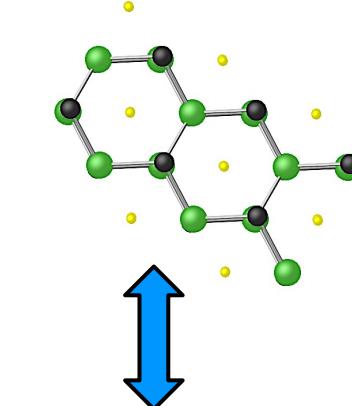
Step 1: materials generation

material generation



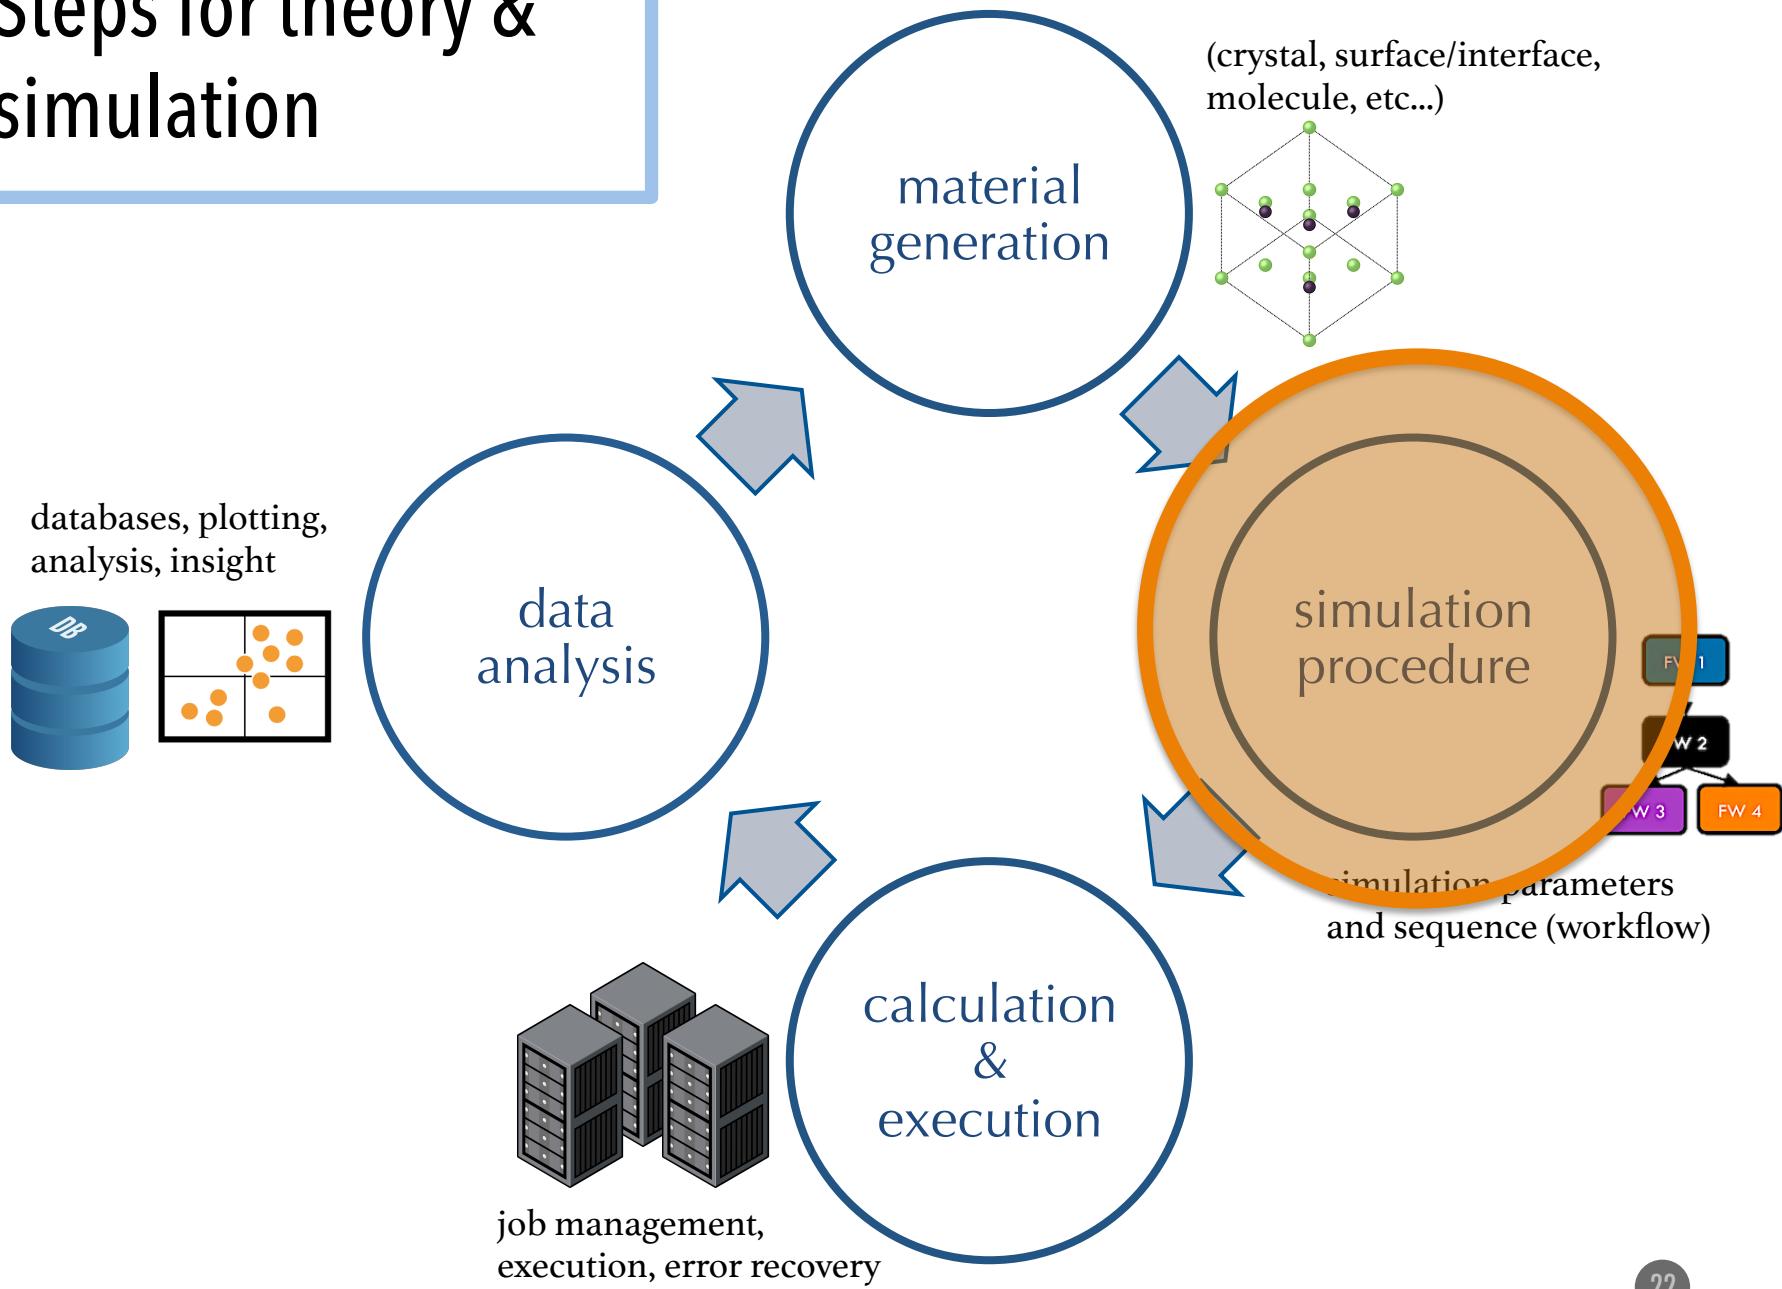
pymatgen

- In yesterday's tutorial, we already saw how pymatgen could help generate models for:
 - crystal structures
 - molecules
 - systems (surfaces, interfaces, etc.)
- Tools include:
 - order-disorder (shown at right)
 - interstitial finding
 - surface / slab generation
 - get structures from Materials Project
- Won't cover materials generation in this presentation!



Example: Order-disorder
resolve partial or mixed
occupancies into a fully
ordered crystal structure
(e.g., mixed oxide-fluoride site
into separate oxygen/fluorine)

Steps for theory & simulation



Step 2: simulation procedure

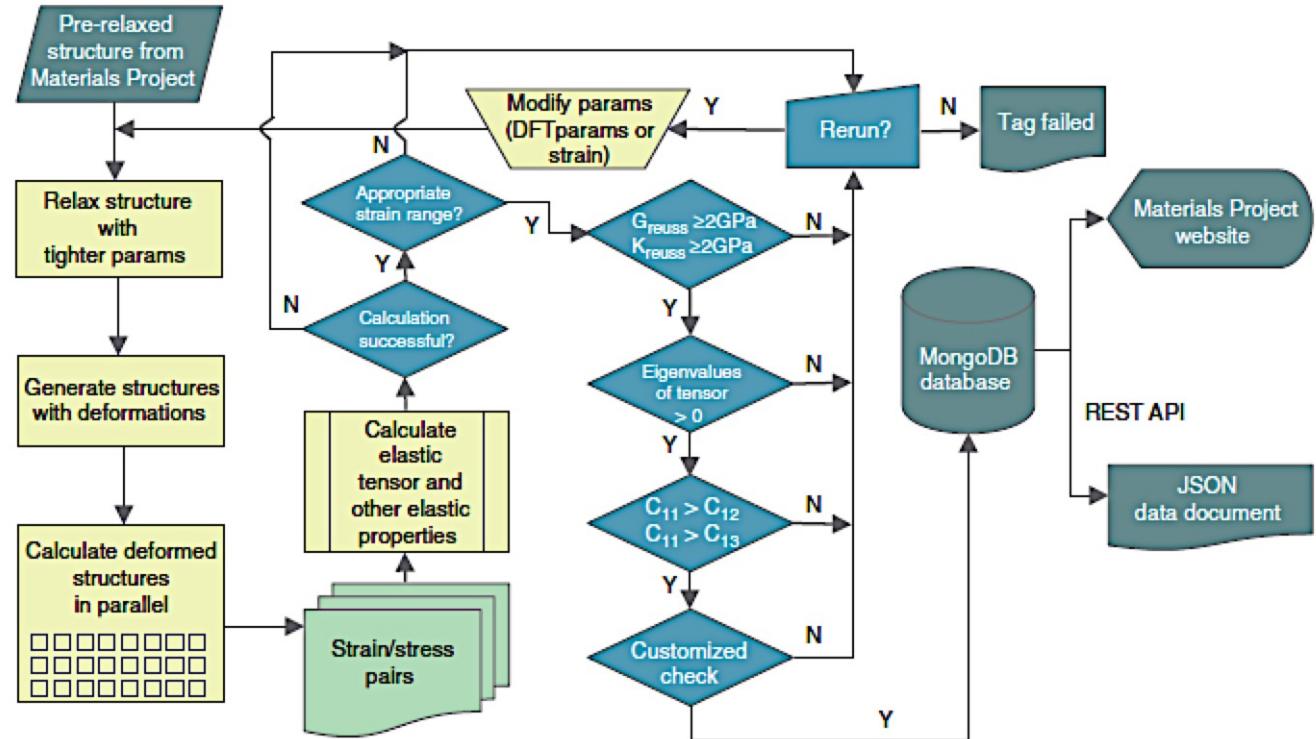
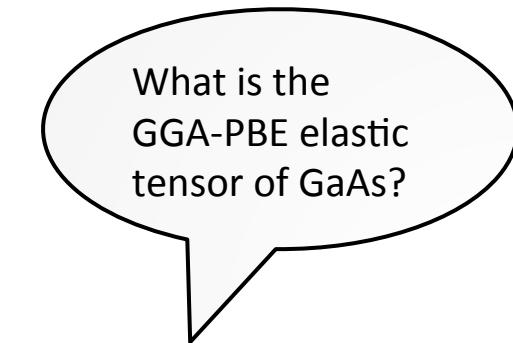


- Simulation procedure requires knowing:
 - i. what sequence of calculations do I need to perform to compute my desired materials property?
 - ii. for each calculation, what is a good set of parameters to use (e.g., functional, energy cutoff, k-point density, etc.)
- Item (i) is largely handled by atomate
- Item (ii) is largely handled by pymatgen, e.g. in the VasplInputSet classes and won't be covered here - although atomate also contains some of this information
- Overall, we are trying to standardize and automate simulation procedures as much as possible

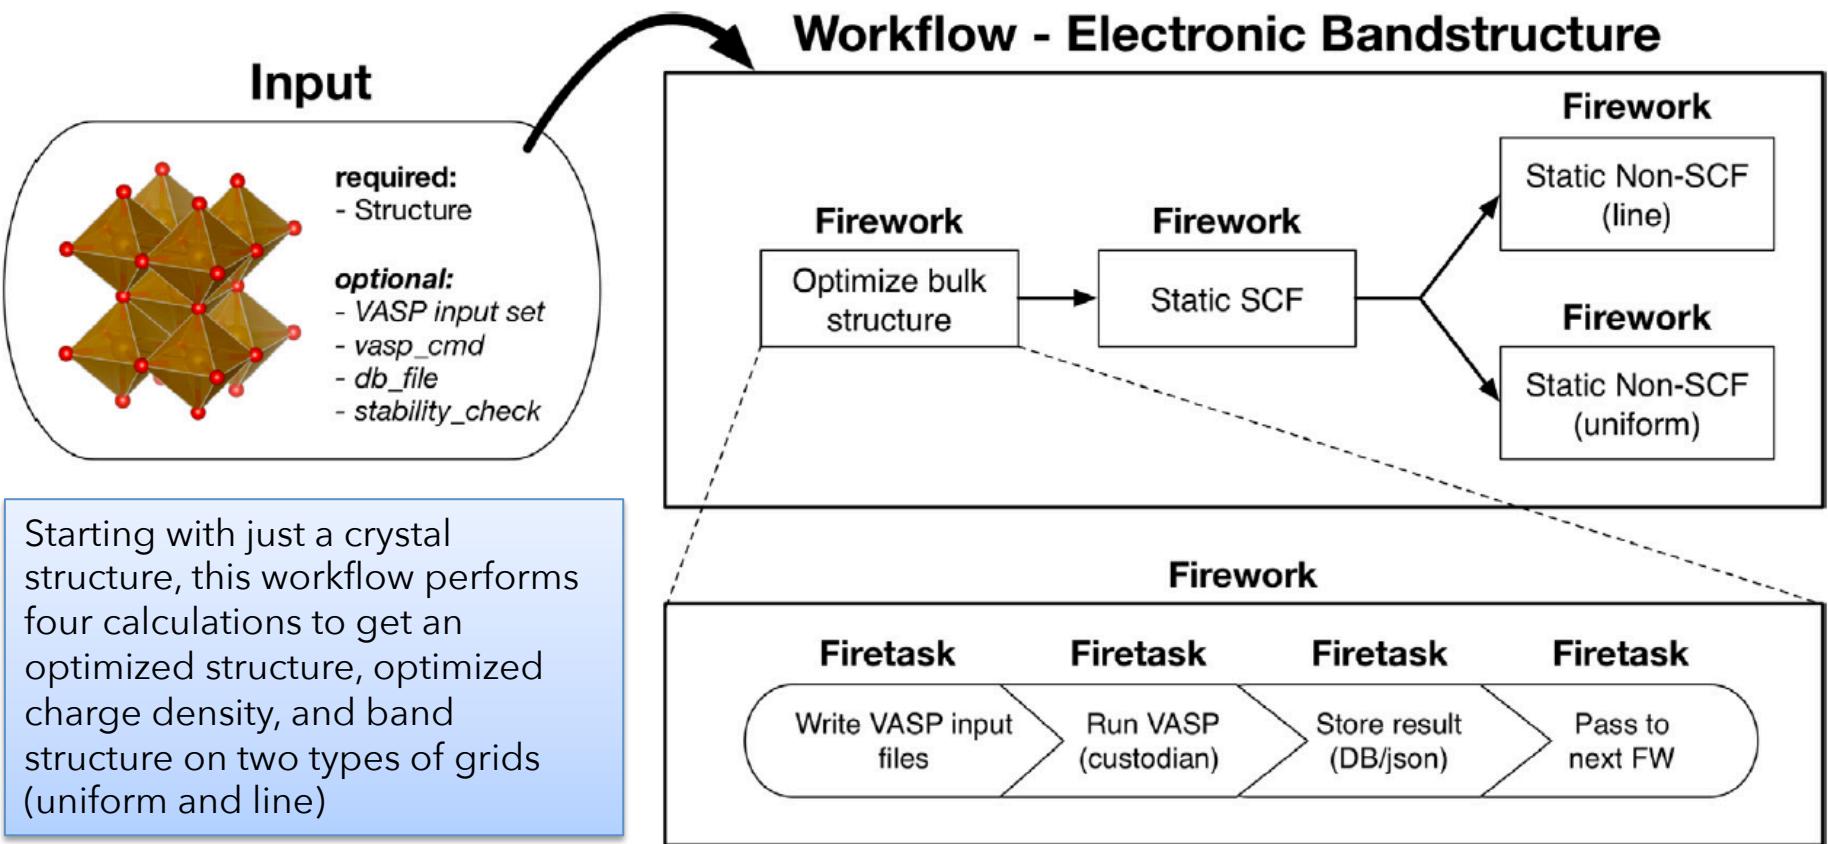
Atomate knows the sequence of calculations needed to compute many kinds of materials properties

atomate 

quickly and automatically translate PI-style (minimal) specifications into well-defined FireWorks workflows



Each simulation procedure in atomate is composed of multiple levels of detail / abstraction



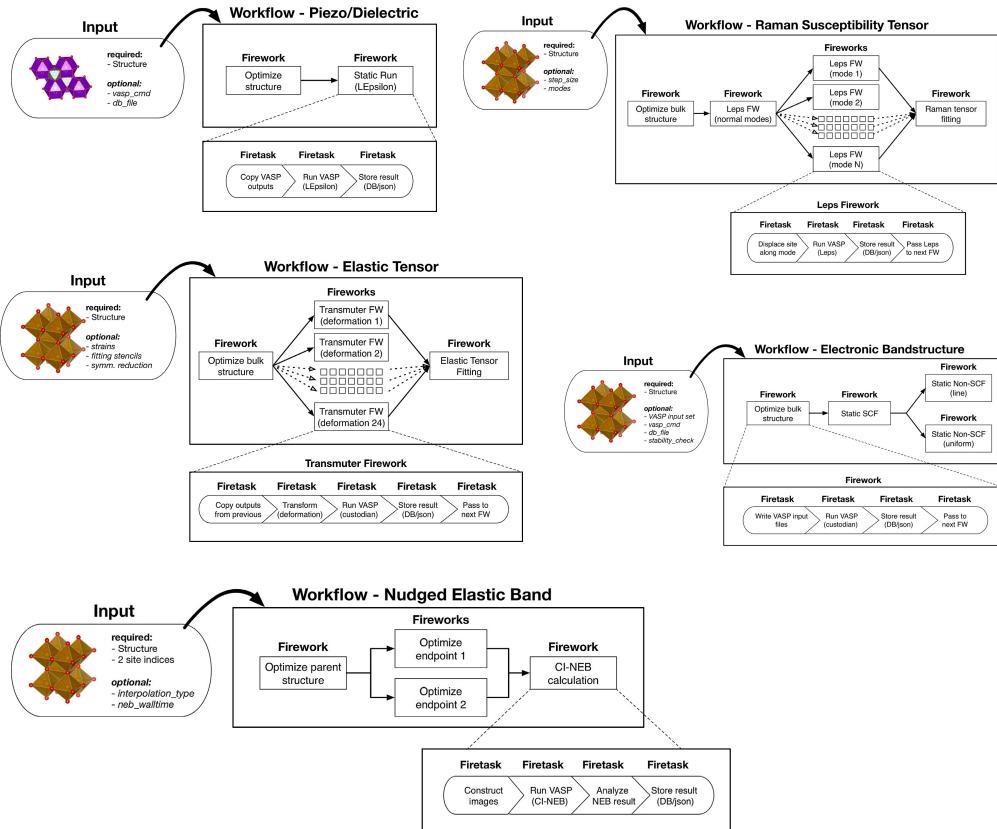
Workflow – complete set of calculations to get a materials property

Firework – one step in the Workflow (typically one DFT calculation)

Firetask – one step in a Firework

Many types of simulation procedures are already available!

Workflows currently available in atomate!



- band structure
- BoltzTraP transport
- spin-orbit coupling
- hybrid functional calcs
- elastic tensor
- piezoelectric tensor
- Raman spectra
- NEB
- GIBBS method
- QH thermal expansion
- AIMD
- FEFF method
- LAMMPS MD

atomate allows you to leverage the prior efforts and knowledge of many researchers

atomate 



All past and present knowledge, from everyone in the group,
everyone previously in the group, and our collaborators,
about how to run calculations



K. Mathew



J. Montoya



S. Dwaraknath



A. Faghaninia



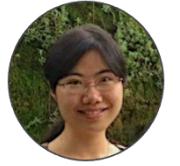
B. Bocklund



T. Smidt



M. Aykol



H. Tang



I.H. Chu



M. Horton



J. Dagdalen



B. Wood



Z.K. Liu



J. Neaton



S.P. Ong

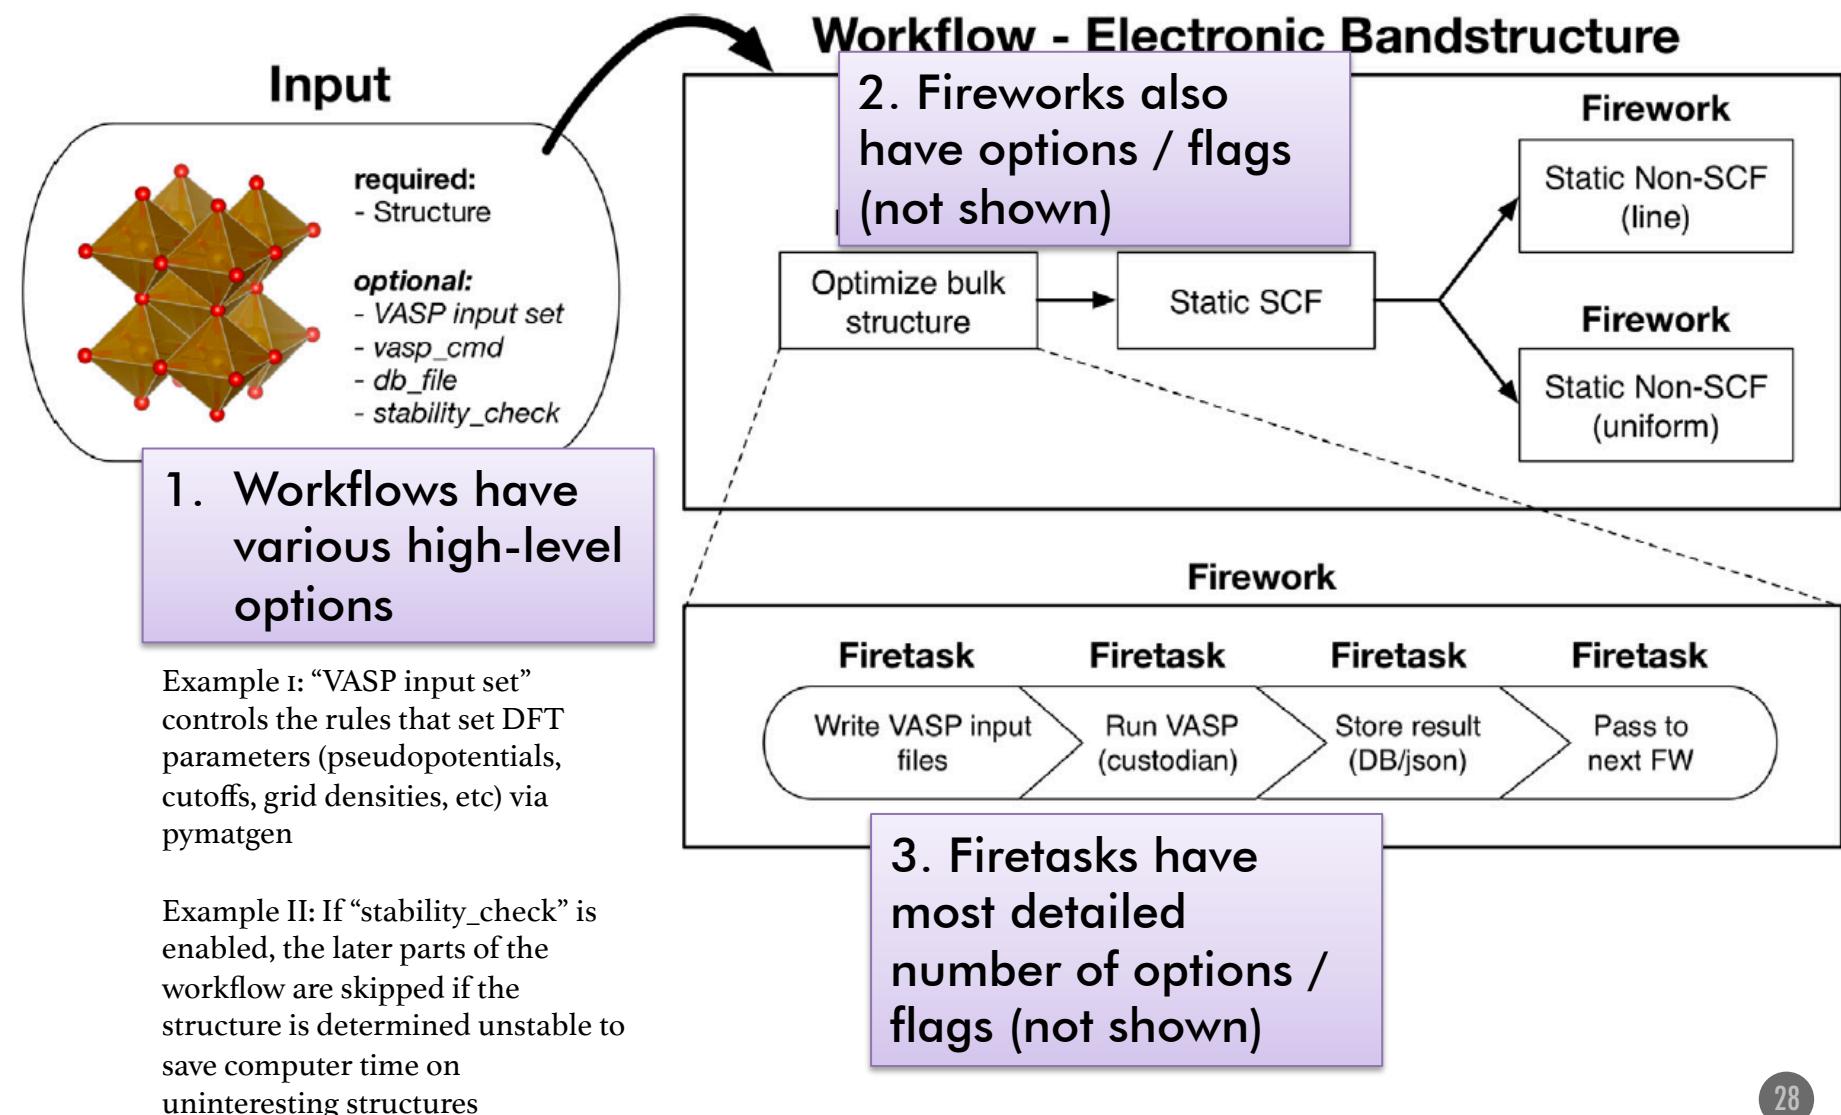


K. Persson



A. Jain

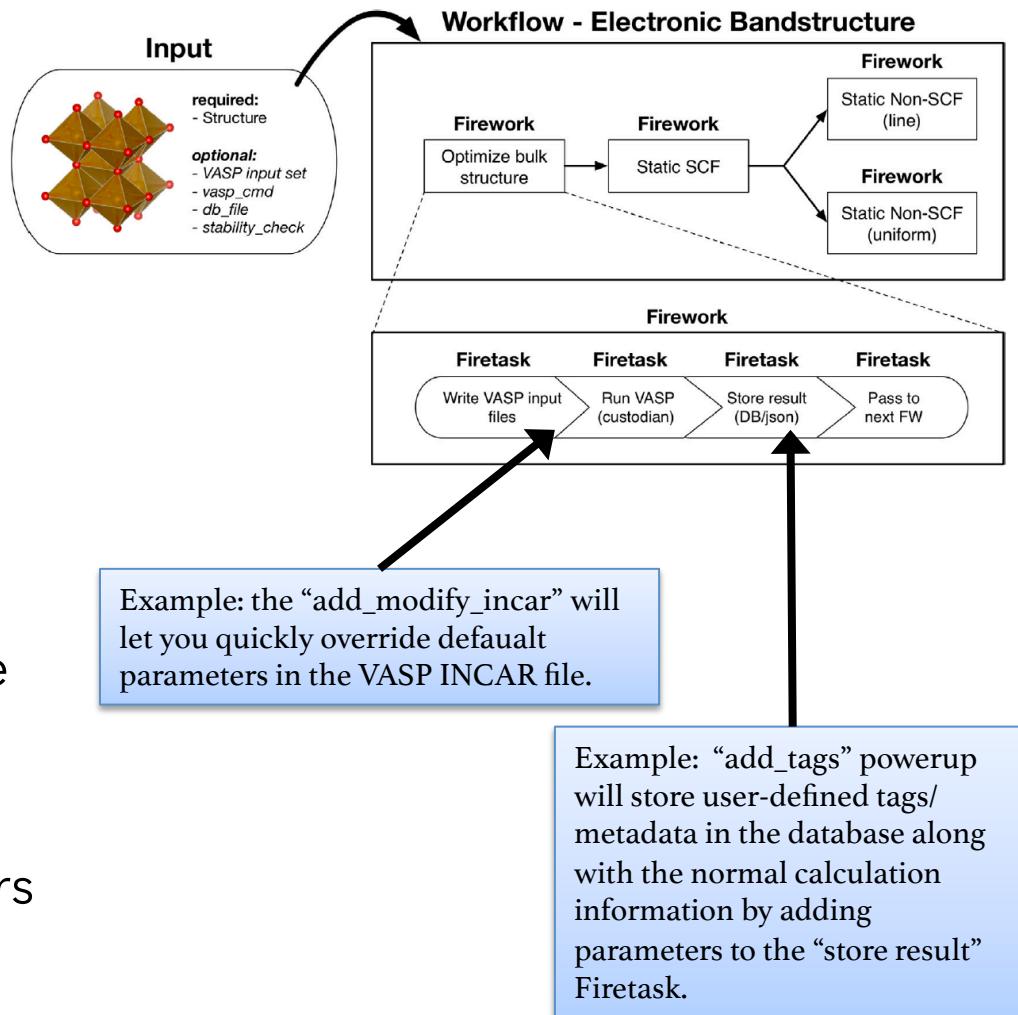
Workflow parameters can be customized at multiple levels of detail



Example II: If “`stability_check`” is enabled, the later parts of the workflow are skipped if the structure is determined unstable to save computer time on uninteresting structures

"Powerups" can also add special features to a workflow, resulting in customization

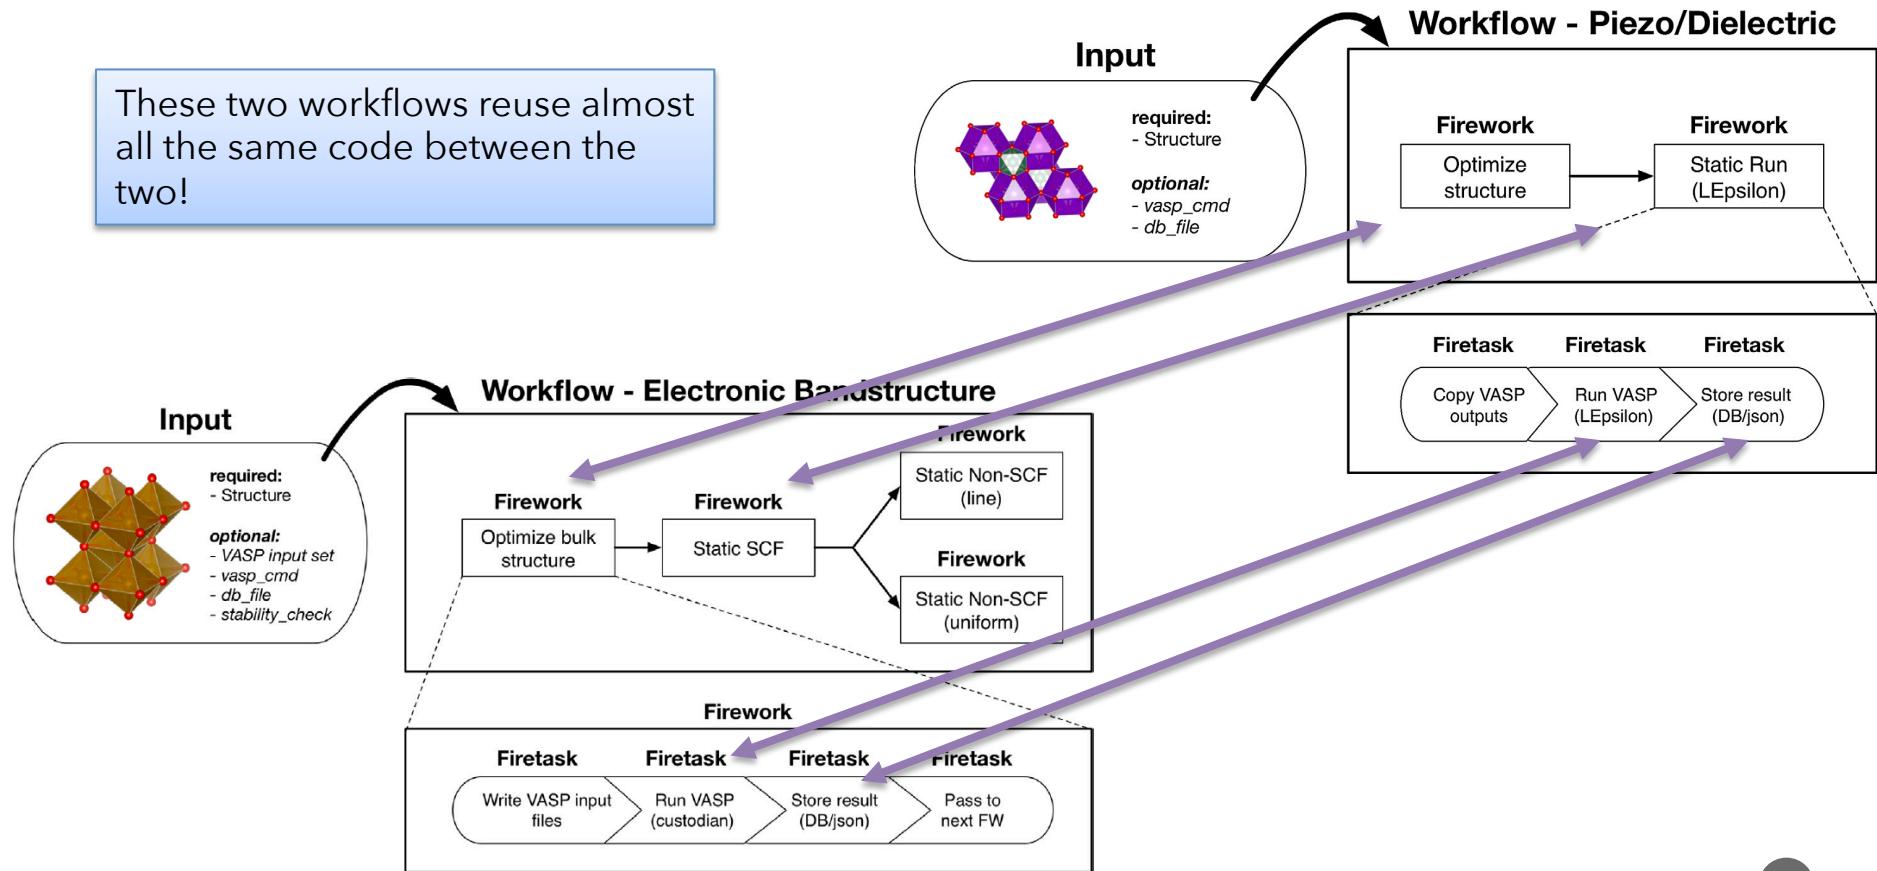
- Atomate's philosophy: base workflow implementations should be clean
 - they only contain the necessary steps to perform a simulation
- What about user preferences like adding tags to certain compounds or tracking the progress of output files in the database?
- "Powerups" act like decorators
 - they take in a workflow and return a workflow with modified properties



You can build workflows from scratch or reuse components to assemble workflows

Multiple workflows are built with the same components stacked together in different ways like Legos

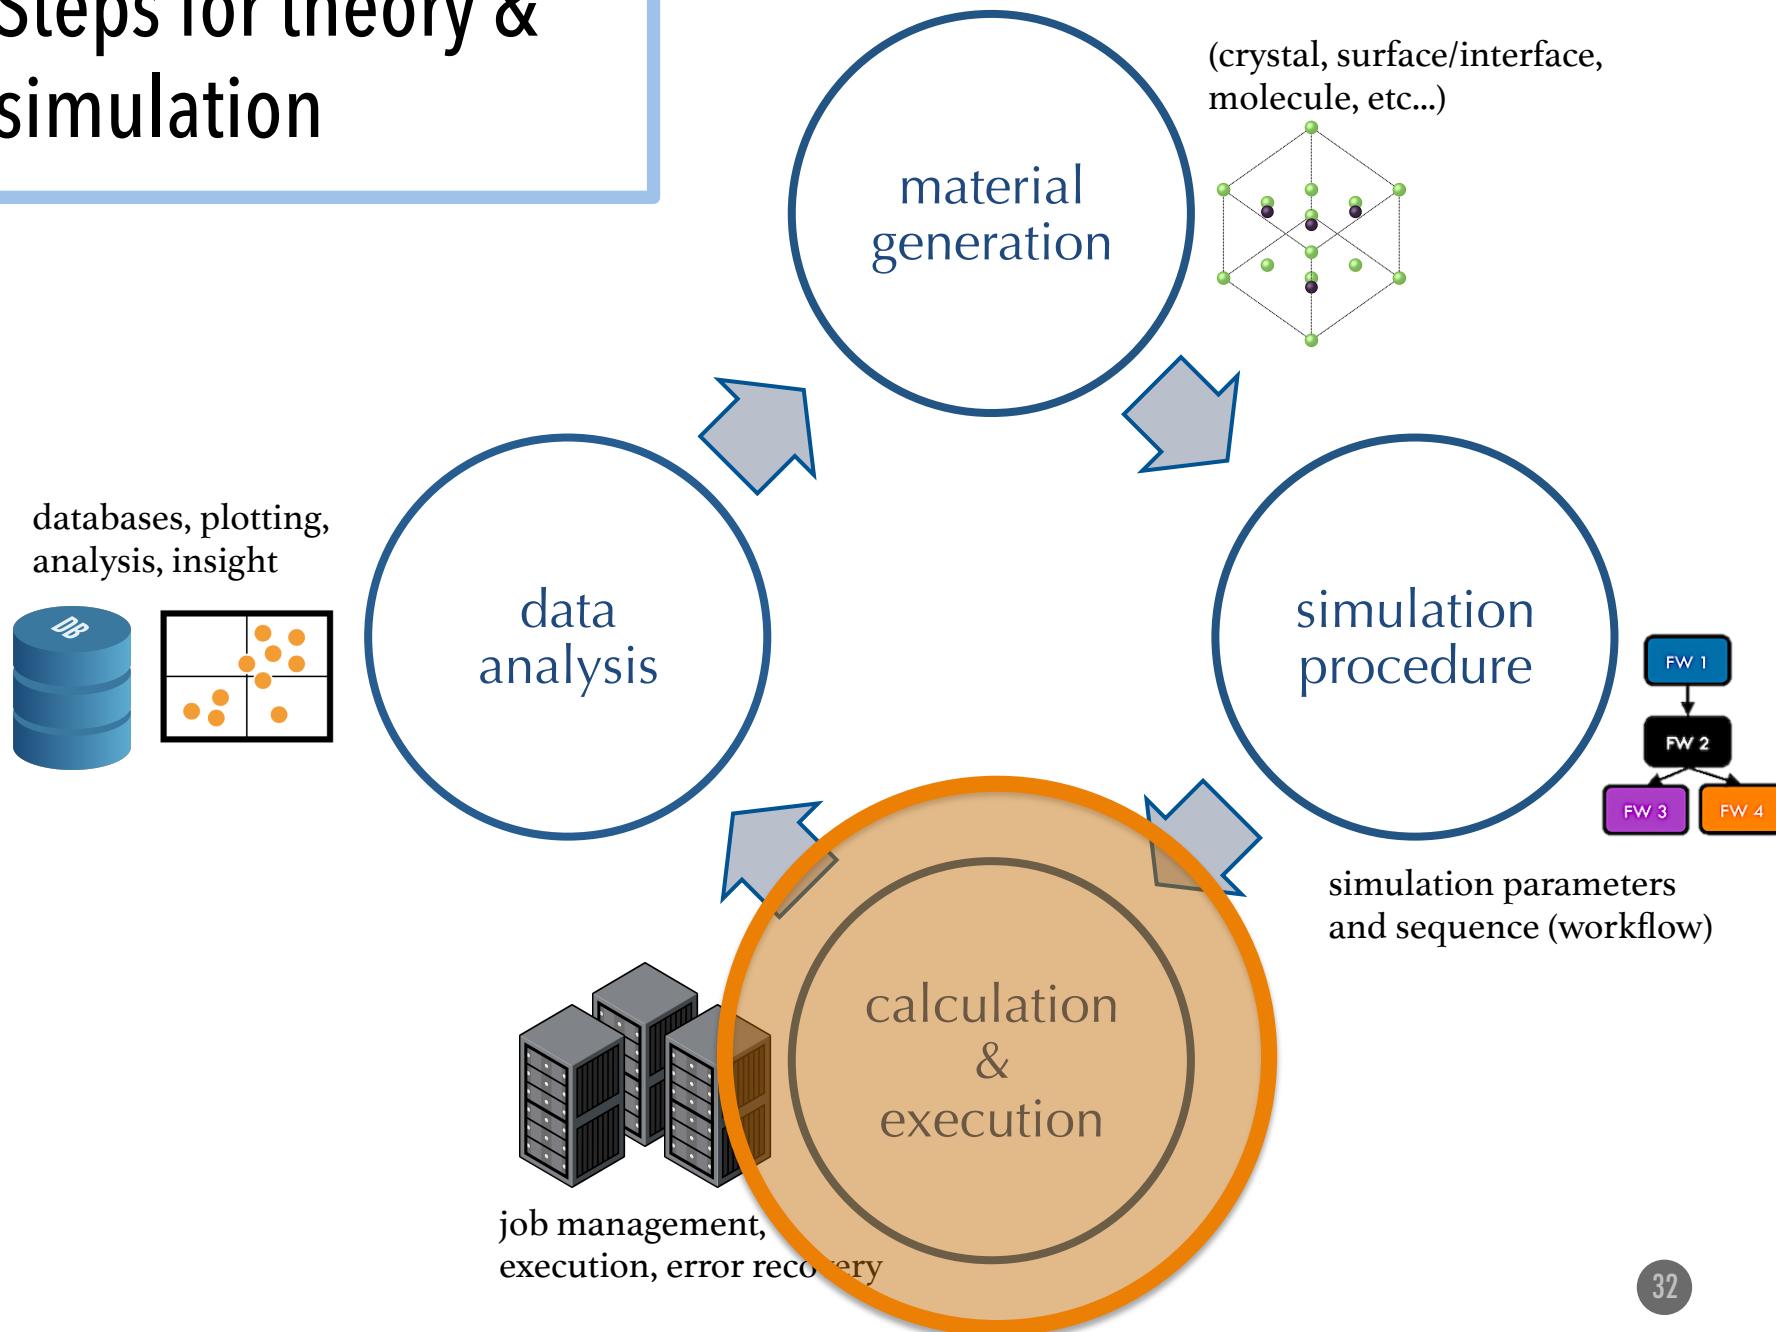
These two workflows reuse almost all the same code between the two!



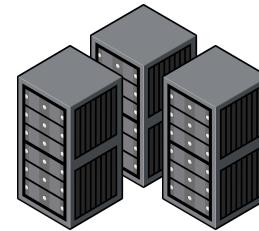
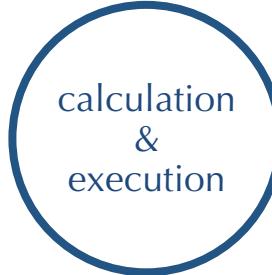
Atomate simulation procedures: the main ideas

- Atomate allows you to start with a material (e.g., crystal structure) and get a full workflow of calculations needed to compute many types of materials properties
- Atomate tries to guess sensible defaults in terms of calculation parameters, but you can always customize these parameters using:
 - Workflow options
 - Firework options
 - Firetask options
 - powerups
 - going back to pymatgen library

Steps for theory & simulation



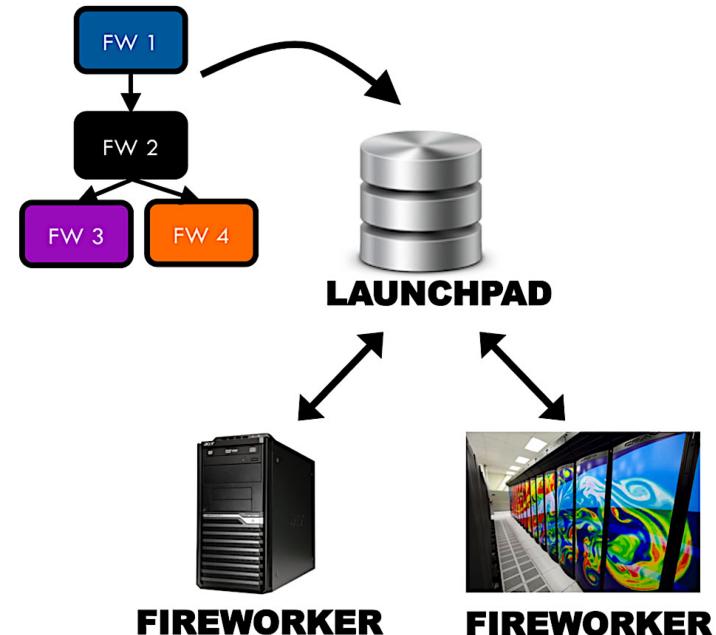
Step 3: calculation & execution



- Once you have the material and the simulation procedure (Workflow), you need to actually execute the workflow on your computing resource
- This includes tasks like:
 - submission to calculation queues
 - customization of any computing-specific parameters
 - e.g., path to VASP executable, number of CPUs to parallelize over
 - recovering from failures / job resubmission
 - coordinating jobs across computing centers
 - managing location of jobs
 - tracking the progress of jobs
- Almost all of this is handled by FireWorks (custodian is used for encoding job recover procedures)

FireWorks - scientific workflow software

- Once you have a Workflow, FireWorks will execute it for you
- FireWorks is an open-source scientific workflow software
- Materials Project, JCESR, and other projects manage their runs with FireWorks
 - >1 million jobs and >100 million CPU-hours for internal projects alone
 - multiple computing clusters
- You can write any kind of workflow
 - e.g., FireWorks is used for graphics processing, machine learning, document processing, and protein folding
 - #1 Google hit for “Python workflow software”, top 5 for general scientific workflow software
- Detailed tutorials are available that cover all the details



FireWorks allows you to write your workflow once and execute (almost) anywhere



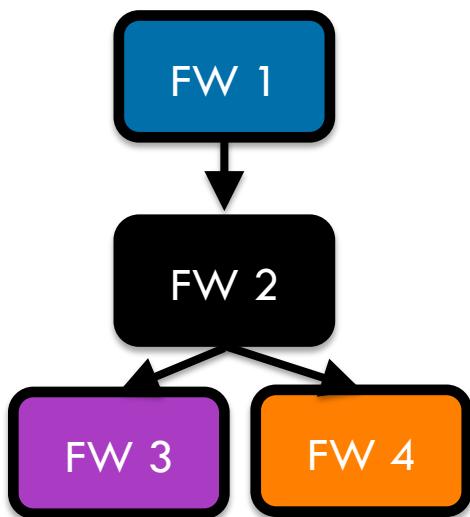
- Execute workflows locally or at a supercomputing center
- Queue systems supported
 - PBS
 - SGE
 - SLURM
 - IBM LoadLeveler
 - NEWT (a REST-based API at NERSC)
 - Cobalt (Argonne LCF)

Job execution (animation)



Directory 1

Directory 2



Dashboard with status of all jobs

 FireWorks

Workflow Dashboard

Newest Workflows		ID: 1573745
B4 C1	READY	
B4_C1--Controller_add_Electronic_Structure_v2		
B4_C1--VASP_db_insertion		
B4_C1--GGA_optimize_structure_(2x)		
B4_C1--Add_to_SNL_database		

Current Database Status		
	Fireworks	Workflows
RUNNING	994	4,728
ARCHIVED	109,576	22,992
WAITING	167,134	0
FIZZLED	76,949	44,682
READY	24,199	18,197
RESERVED	799	0
COMPLETED	1,187,041	111,402
DEFUSED	6,588	3,656
TOTAL	1,573,280	205,657

Newest Workflows		ID: 1573739
Ba2 Fe1 Nb1 O6	READY	
Ba2_Fe1_Nb1_O6--Controller_add_Electronic_Structure_v2		
Ba2_Fe1_Nb1_O6--VASP_db_insertion		
Ba2_Fe1_Nb1_O6--GGAU_optimize_structure_(2x)		
Ba2_Fe1_Nb1_O6--VASP_db_insertion		
Ba2_Fe1_Nb1_O6--GGA_optimize_structure_(2x)		
Ba2_Fe1_Nb1_O6--Add_to_SNL_database		

Newest Workflows		ID: 1573733
Ba2 Fe1 Nb1 O6	READY	
Ba2_Fe1_Nb1_O6--Controller_add_Electronic_Structure_v2		
Ba2_Fe1_Nb1_O6--VASP_db_insertion		
Ba2_Fe1_Nb1_O6--GGAU_optimize_structure_(2x)		
Ba2_Fe1_Nb1_O6--VASP_db_insertion		
Ba2_Fe1_Nb1_O6--GGA_optimize_structure_(2x)		
Ba2_Fe1_Nb1_O6--Add_to_SNL_database		

Summary Reports

Get a report of all jobs for the past:

- 30 minutes
- 24 hours
- 7 days
- 30 days
- 6 months
- 24 months
- 10 years

For more reporting options, use the "ipad report --help" command line tool.

 FireWorks

Workflow 1951337

```
graph TD; A[K1_Nb1_O3--Add_to_SNL_database] --> B[K1_Nb1_O3--GGA_optimize_structure_(2x)]; B --> C[K1_Nb1_O3--VASP_db_insertion]; C --> D[K1_Nb1_O3--Controller_add_Electronic_Structure_v2]
```

READY COMPLETED

[Collapse](#) [Expand](#) [Toggle](#) [Toggle level1](#) [Toggle level2](#)

```
{  
    + created_on: "2017-07-28T17:00:13.517000",  
    + launch_dirs: { - },  
    + links: { - },  
    + metadata: { - },  
    name: "K1 Nb1 O3",  
    + parent_links: { - },  
    state: "RUNNING",  
    + states: { - },  
    updated_on: "2017-08-01T08:01:19.128000"  
}
```

© Copyright 2015, FireWorks.

37

Job provenance and automatic metadata storage

what machine
what time
what directory

what was the output

when was it queued

when did it start running

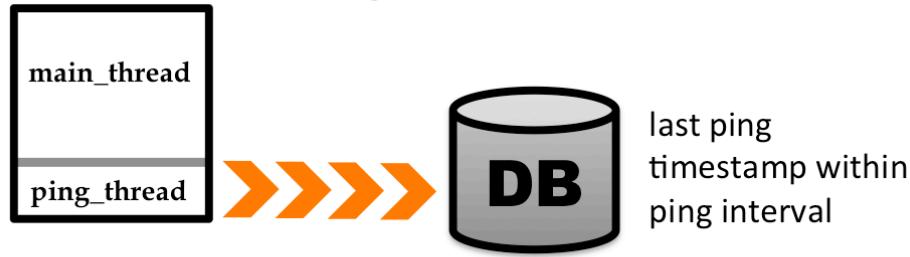
when was it completed

▼ _id	ObjectId("53f4d749835a896053f128c2")	Object id
_id	ObjectId("53f4d749835a896053f128c2")	Object id
► fworker		Object, 4 items
time_start	2014-08-20T23:28:07.901440	String
trackers		Array, no items
ip	10.32.47.192	String
fw_id	952913	Integer
time_end	2014-08-20T23:28:44.838513	String
reservedtime_secs	22462.290943	Double
runtime_secs	36.937073	Double
state	COMPLETED	String
launch_dir	/global/scratch2/sd/matcomp/mp_prod/block_2014-08-16-07-15-18-137142/launcher_2014-08-20-17-13-45-547255	String
host	mc0853	String
launch_id	828022	Integer
▼ action		Object, 7 items
► update_spec		Object, 8 items
mod_spec		Object, 1 item
► stored_data		Object, 1 item
exit	false	Boolean
detours		Array, no items
additions		Array, no items
defuse_children	false	Boolean
▼ state_history		Array, 3 items
▼ 0		Object, 4 items
updated_on	2014-08-20T17:13:45.610503	String
state	RESERVED	String
reservation_id	9919235	String
created_on	2014-08-20T17:13:45.610497	String
▼ 1		Object, 3 items
updated_on	2014-08-20T23:28:44.756385	String
state	RUNNING	String
created_on	2014-08-20T23:28:07.901440	String
▼ 2		Object, 2 items
state	COMPLETED	String
created_on	2014-08-20T23:28:44.838513	String

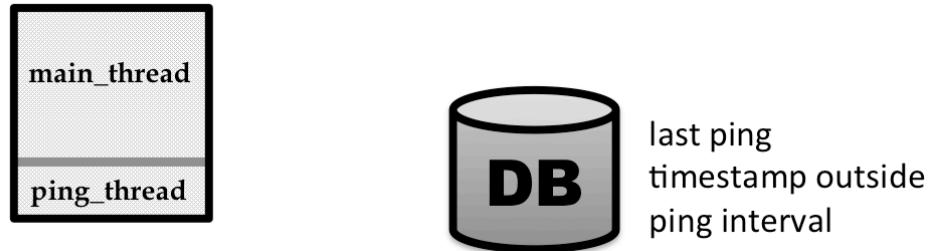
Detect and rerun failures

- All kinds of failures can be detected and rerun
 - Soft failures (job quits with error code)
 - hard failures (computing center goes down)
 - human errors

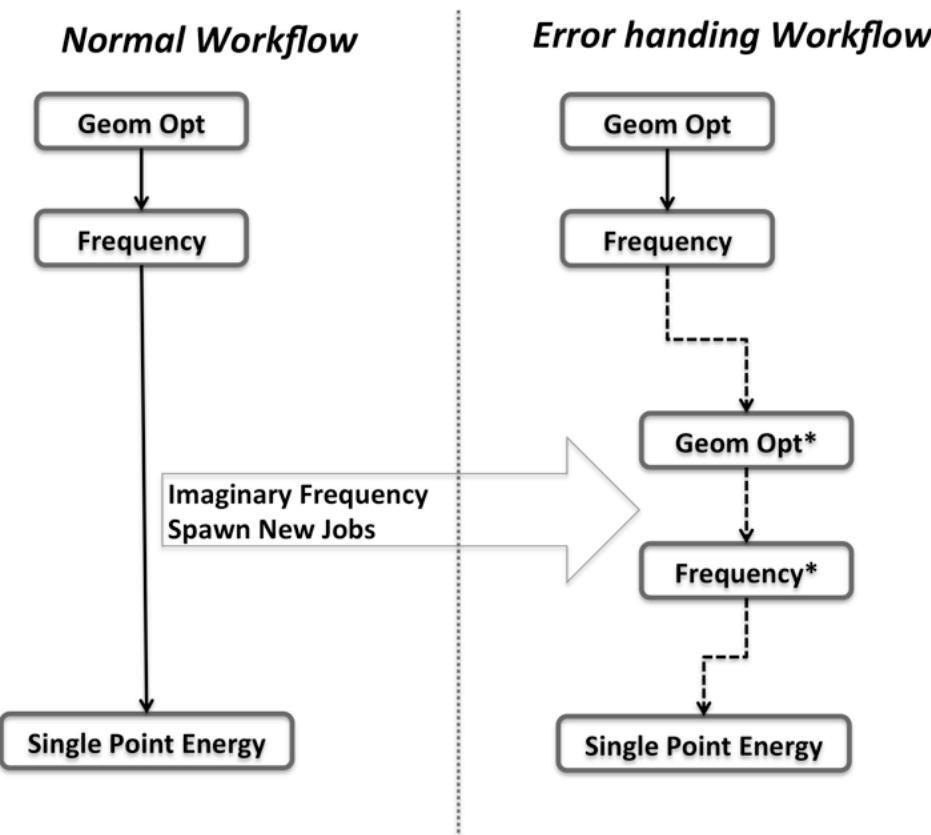
“alive” + running



“dead” job



“Dynamic workflows” let you program intelligent, reactive workflows



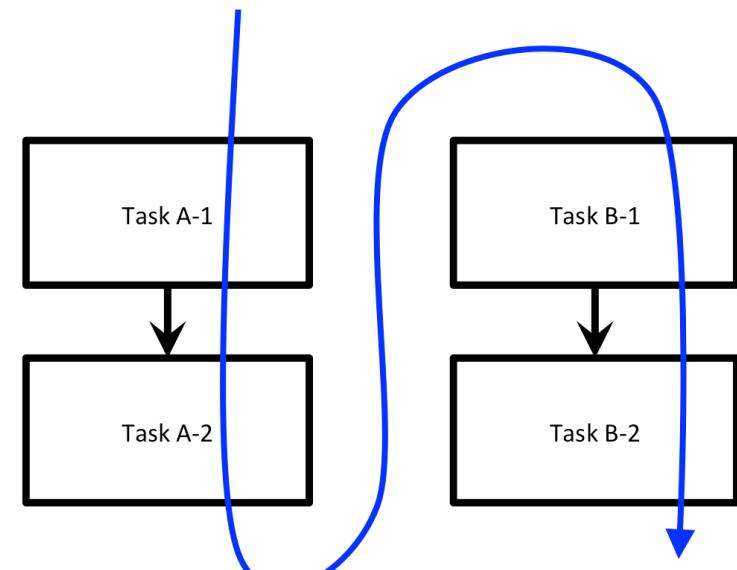
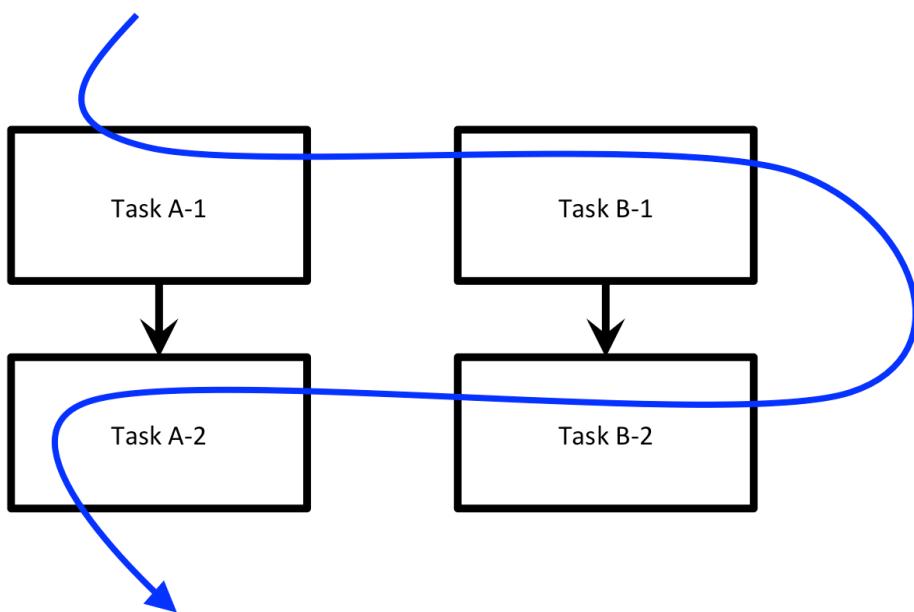
Xiaohui can replace himself with

```
001000100001000001101101100000011010111001100001010  
10000111011000110010010 0111110100110001011001010  
000001011001100101 100111111 1000100001101010010000  
1111010010100011 10111111 01011101010101 001  
010111011101001 00100000010 100000101010111  
0000100001110101 10010100011 01000111 001 1  
0010001001110100 0011011100 1100 000010 0000  
111010111000000011101110100110 000 00010 0000  
1011001111111001000001100 1 1 10 0110  
01001110000111 1 0011001000 100 0110  
010100001 1 1 11101111 01010001 00100  
01111011 1 0 10010 000010  
01 11111 111 0101100101 0 0  
001 11 1 0011110 0111100 11 110  
1 0101000 110001 0000111101 010  
000111000110110100 001100100 11 1 0  
00010110001001100 01111000011110 00 0  
101101011100101000010000010111011 0 1  
000001000100010100001110101110000 1 0 10  
11001000111011011111000111011011 000111000  
000 00010001110100010011110111 1111101 11
```

digital Xiaohui,
programmed into
FireWorks

Customize job priorities

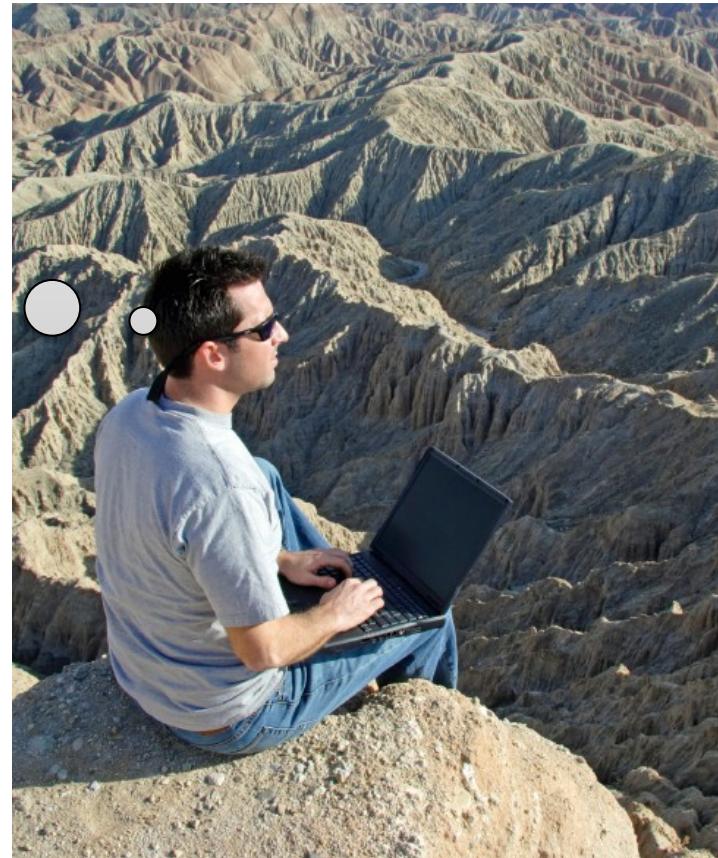
- Within workflow, or between workflows
- Completely flexible and can be modified / updated whenever you want



Track output files remotely

- Can bring up the last few lines of each of your output files – and can be combined with queries / filters

Now seems like a good time to bring up the last few lines of the OUTCAR of all failed jobs...

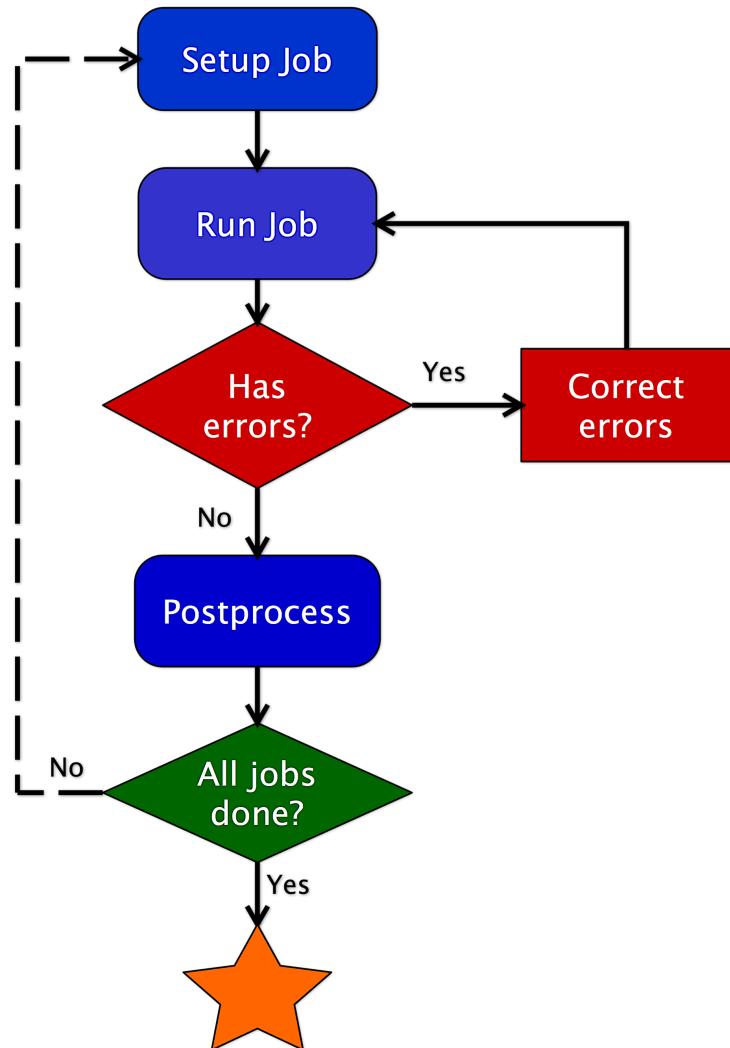


FireWorks workflow software: the main ideas

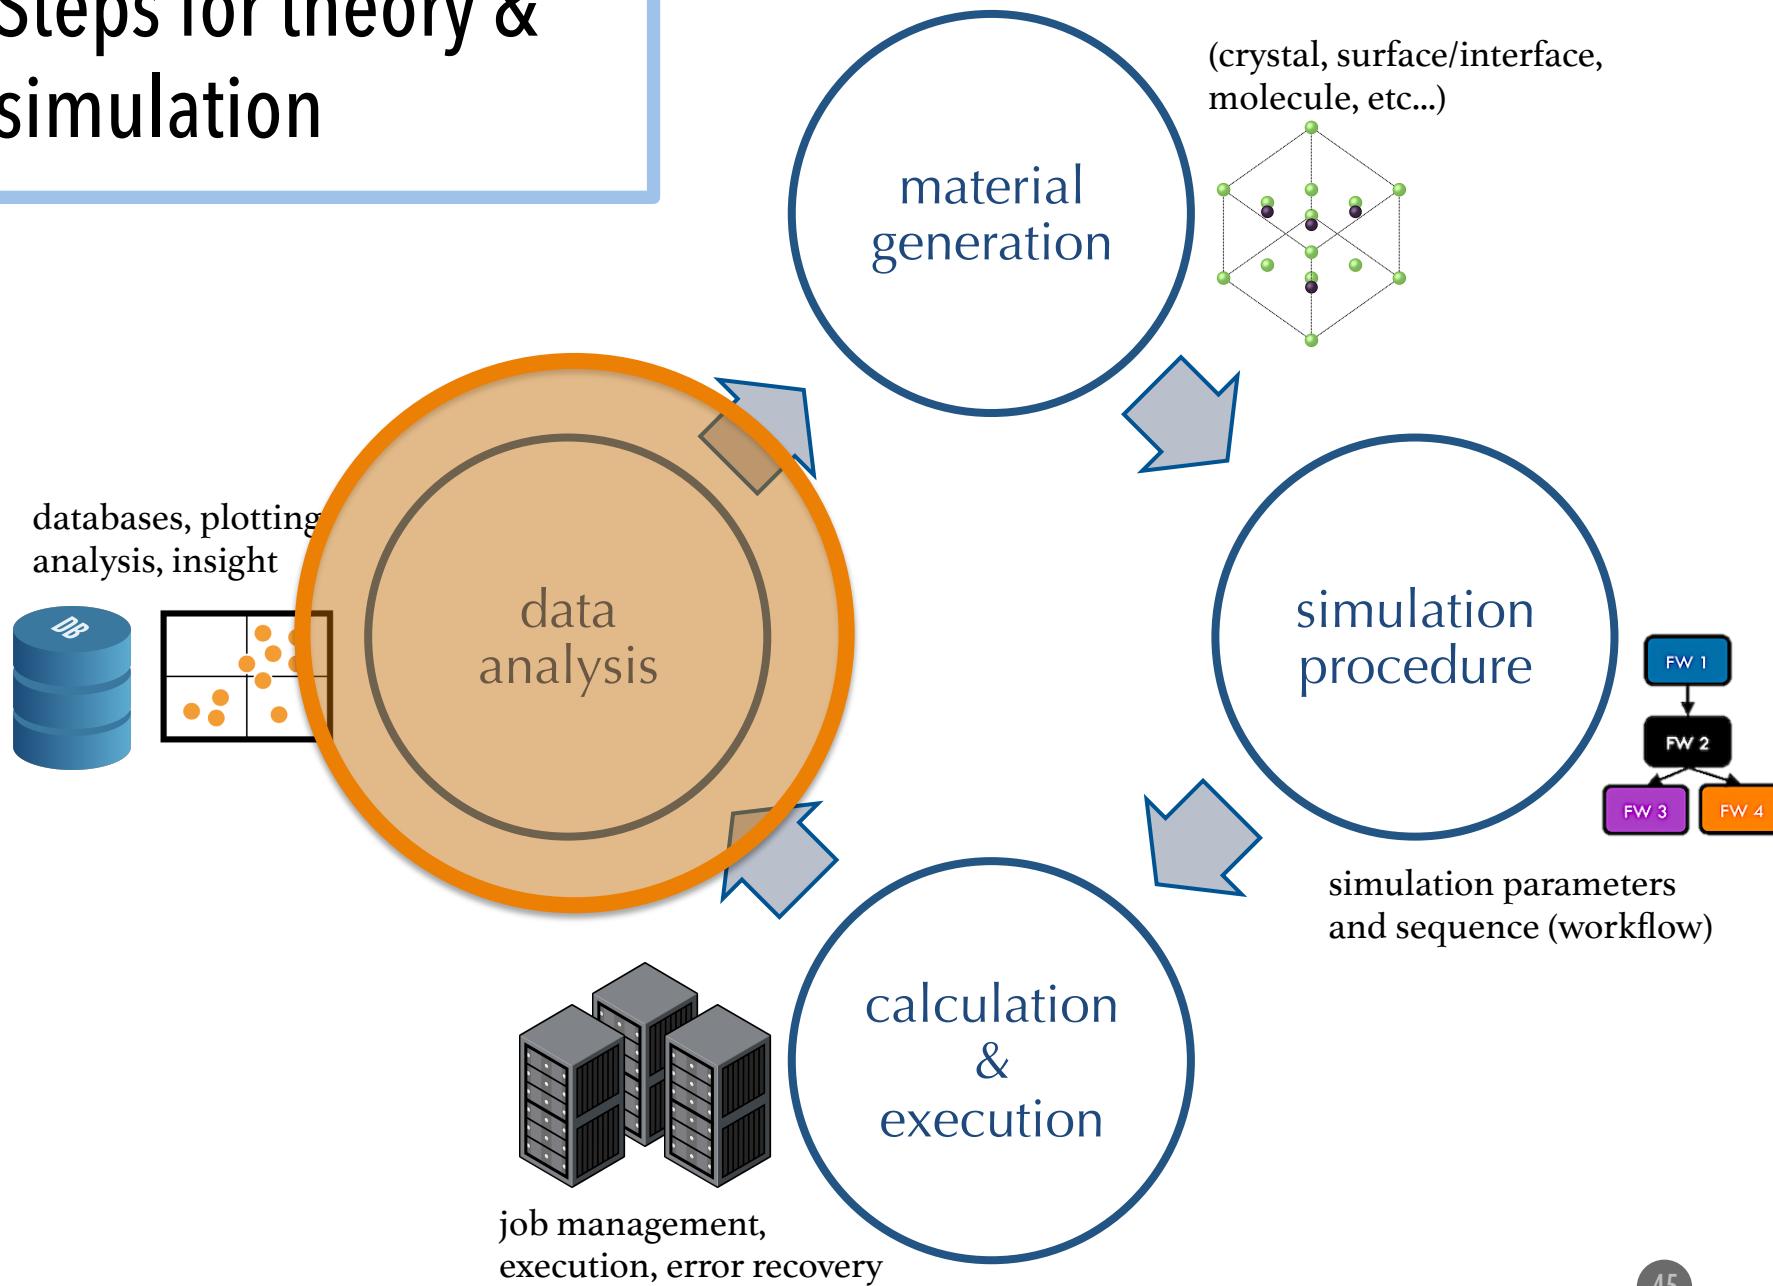
- FireWorks is used to execute Workflow objects at supercomputing centers
- It is very well-suited to high-throughput applications and has been used to execute millions of jobs and is well tested
- There are many features and advantages to using FireWorks as your workflow manager, which has made it one of the most popular scientific workflow software in use today

An aside on "custodian": instead of running VASP, have custodian run VASP on your behalf

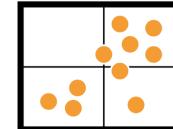
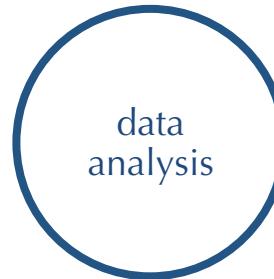
- Custodian can wrap around an executable (e.g., VASP)
 - i.e., run custodian instead of directly running VASP
- During execution, custodian will monitor output files and detect errors / problems
 - e.g., if ZPOTRF error detected, rerun with ISYM=0
 - ever-expanding library of fixes (currently over 30 fixes for VASP!)
- There is more you can do with custodian (e.g., simple workflows), but it won't be covered here
- Currently has fixes for: VASP, Q-Chem, NWChem



Steps for theory & simulation



Step 4: data analysis

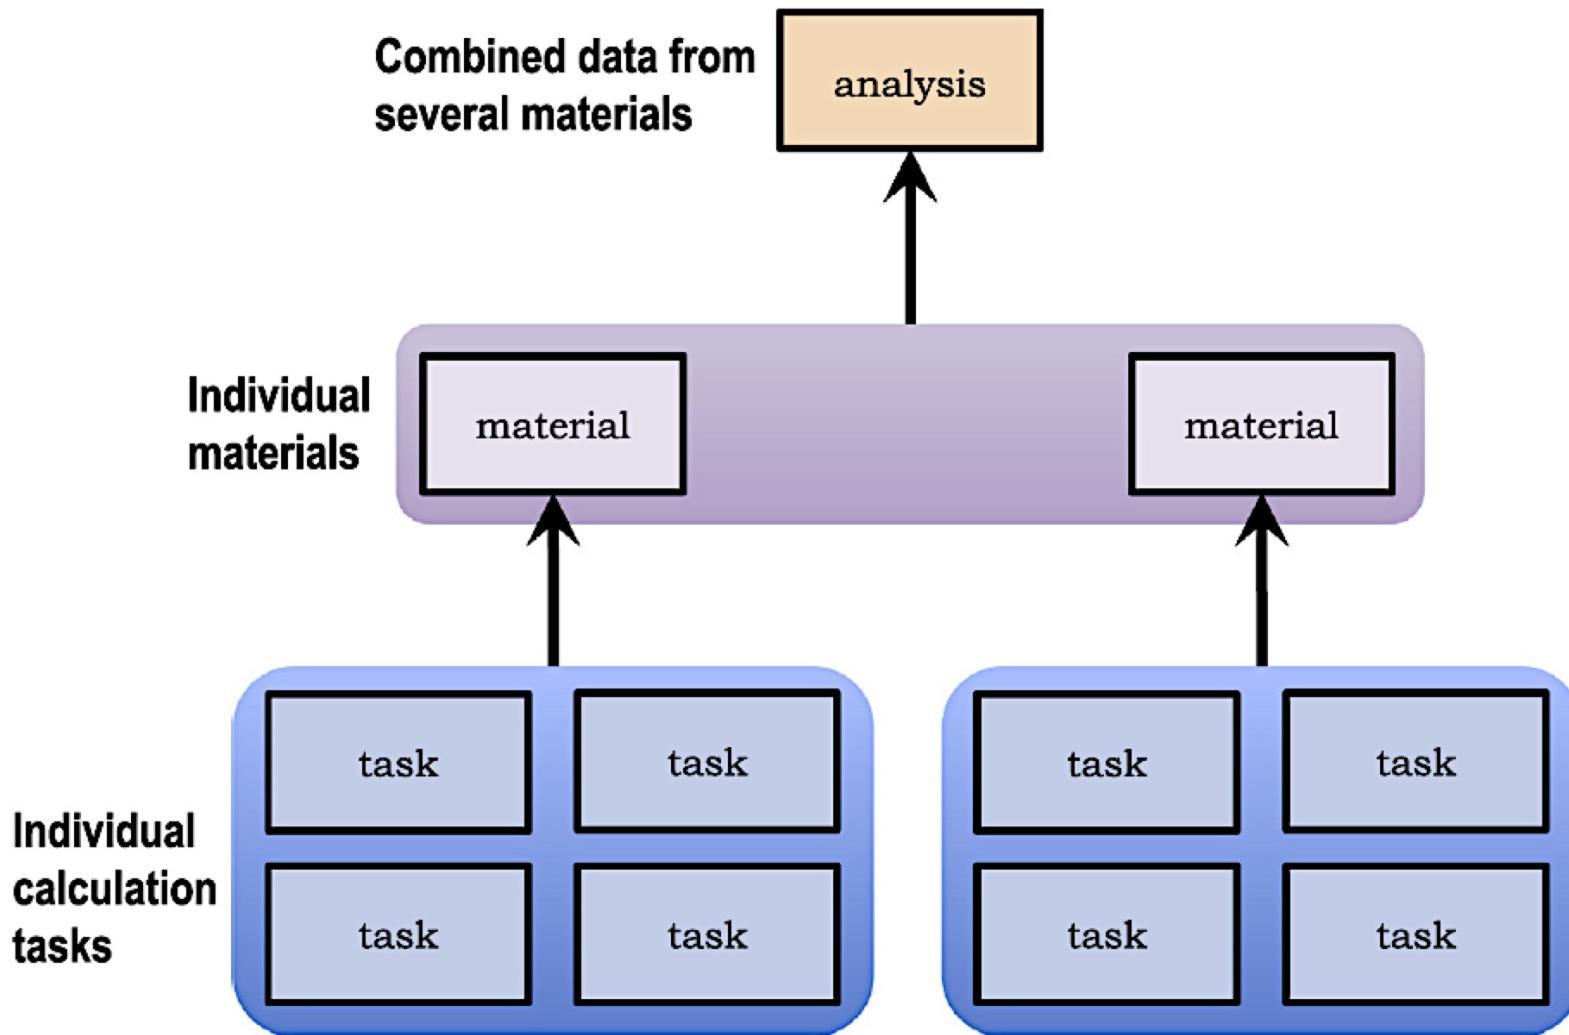


pymatgen
atomate
matminer

- Data analysis can involve:
 - i. Creating databases of materials and their properties for easy searching and data retrieval
 - ii. Conventional scientific analyses and scientific plotting (e.g., generating phase diagrams, plotting band structures)
 - iii. Data mining methods
- Item (i) is largely handled by atomate
- Item (ii) is largely handled by pymatgen
- Item (iii) is largely handled by matminer

Atomate - builders framework

"Builders" start with base collections in a database and create higher-level collections that summarize information or add metadata



What kinds of builders are available in atomate?

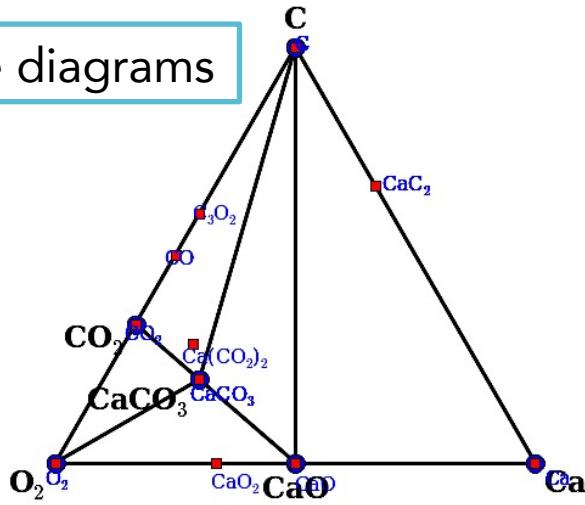
- “Materials builder” - combine all calculations (“tasks”) on a single crystal (as determined automatically by structure matching algorithms) into a single document
 - e.g., combine data from band structure, dielectric, and elastic constant workflows performed on the same material into a single report on that material that contains all information
- Add new information based on algorithmic analysis to the database, e.g. structure dimensionality or phase stability
- Merge external information such as user-defined tags, experimental data (from a spreadsheet), or data from Materials Project to the computational report of a material
- If you are using atomate without using builders, you are missing out on one of the most useful features!

Data analysis - the role of atomate

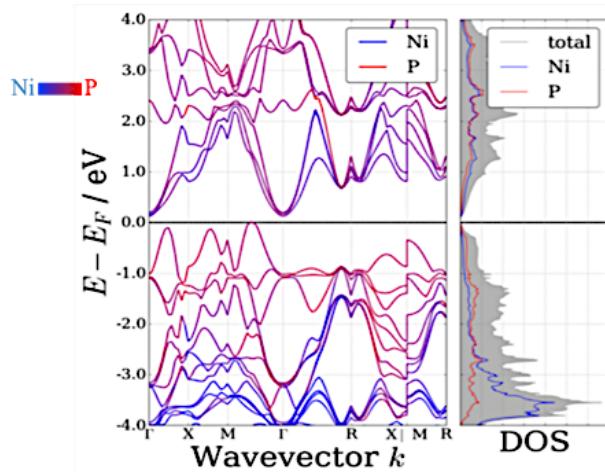
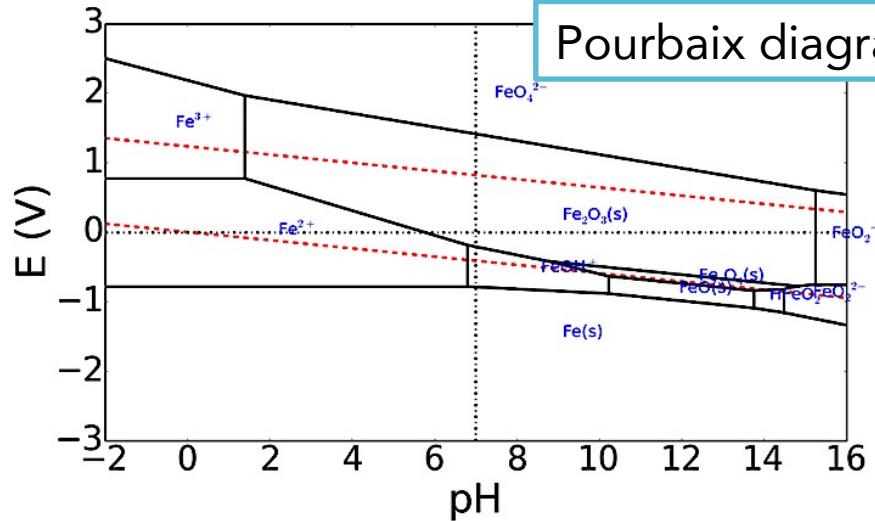
- Atomate helps organize computational data into databases
- However, atomate does not do scientific analysis or plotting of the data - it can just generate and grab the data (often as pymatgen objects)
- Pymatgen is used to analyze the data
 - we won't cover pymatgen capabilities in depth since it is covered elsewhere in the workshop

pymatgen - examples of analyses

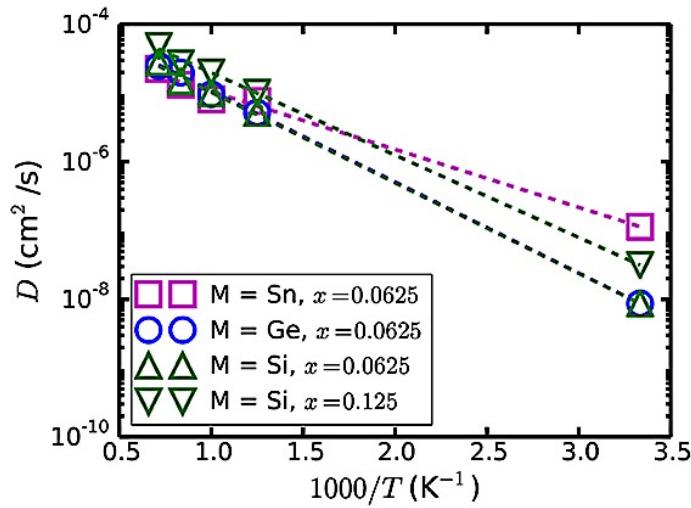
phase diagrams



Pourbaix diagrams

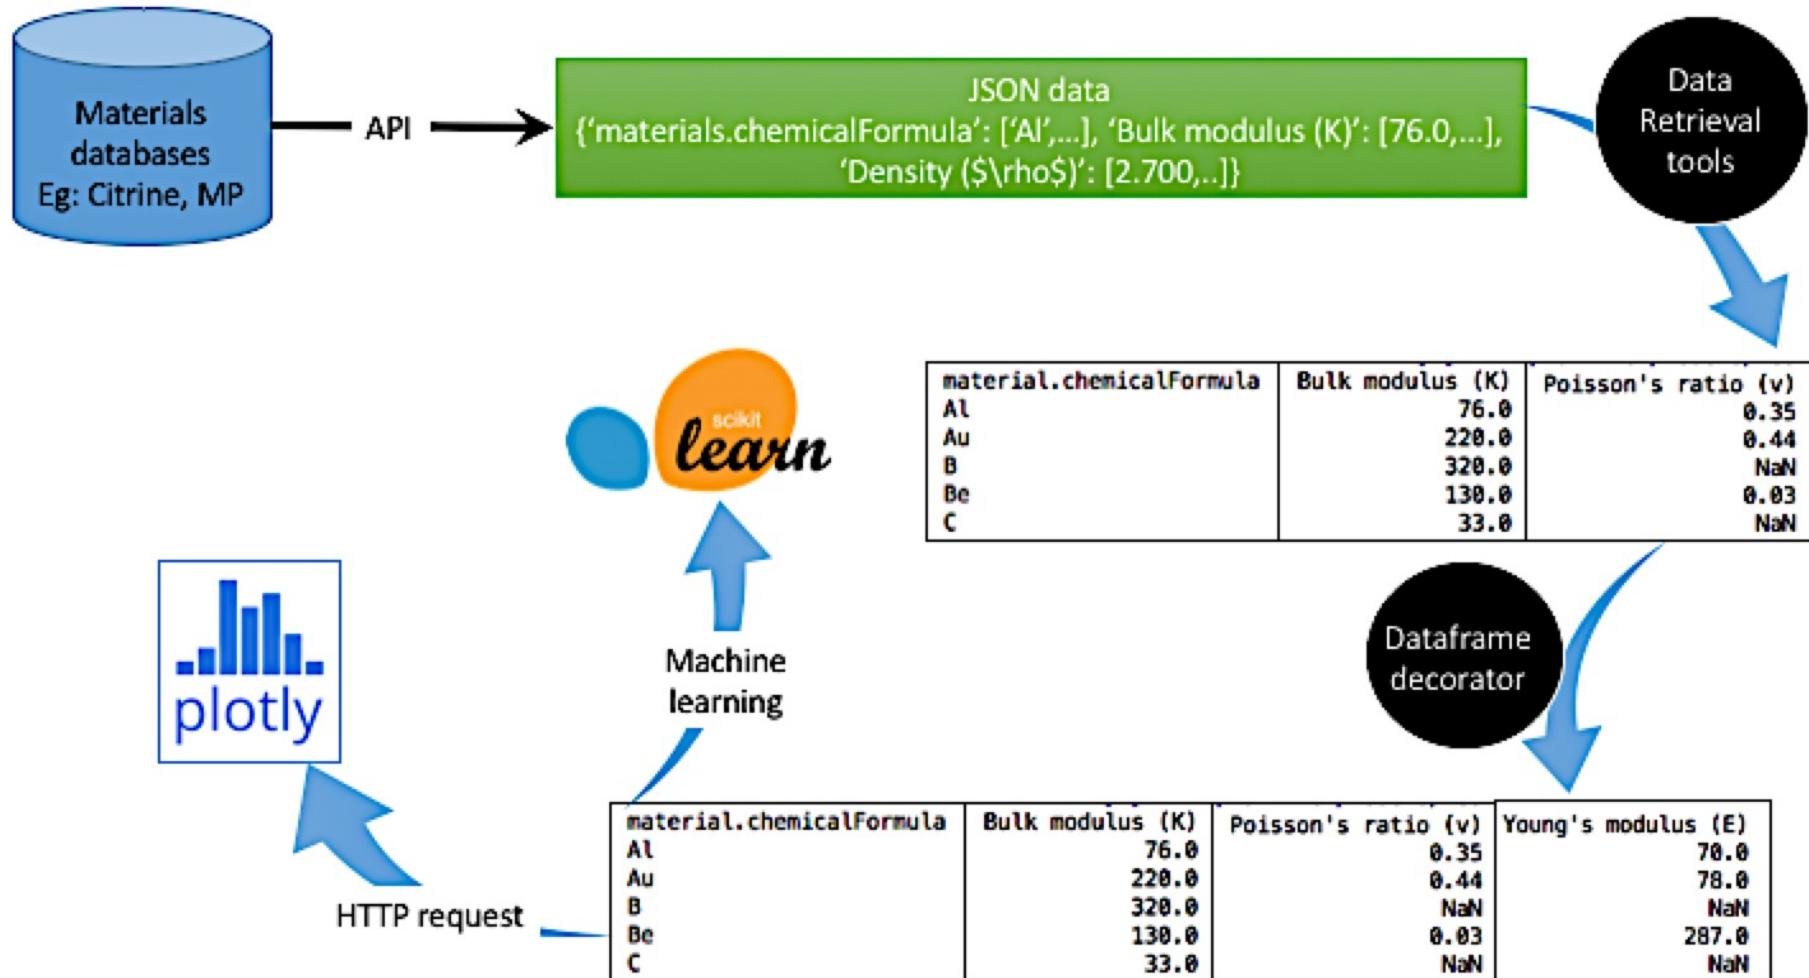


band structure analysis



diffusivity from MD

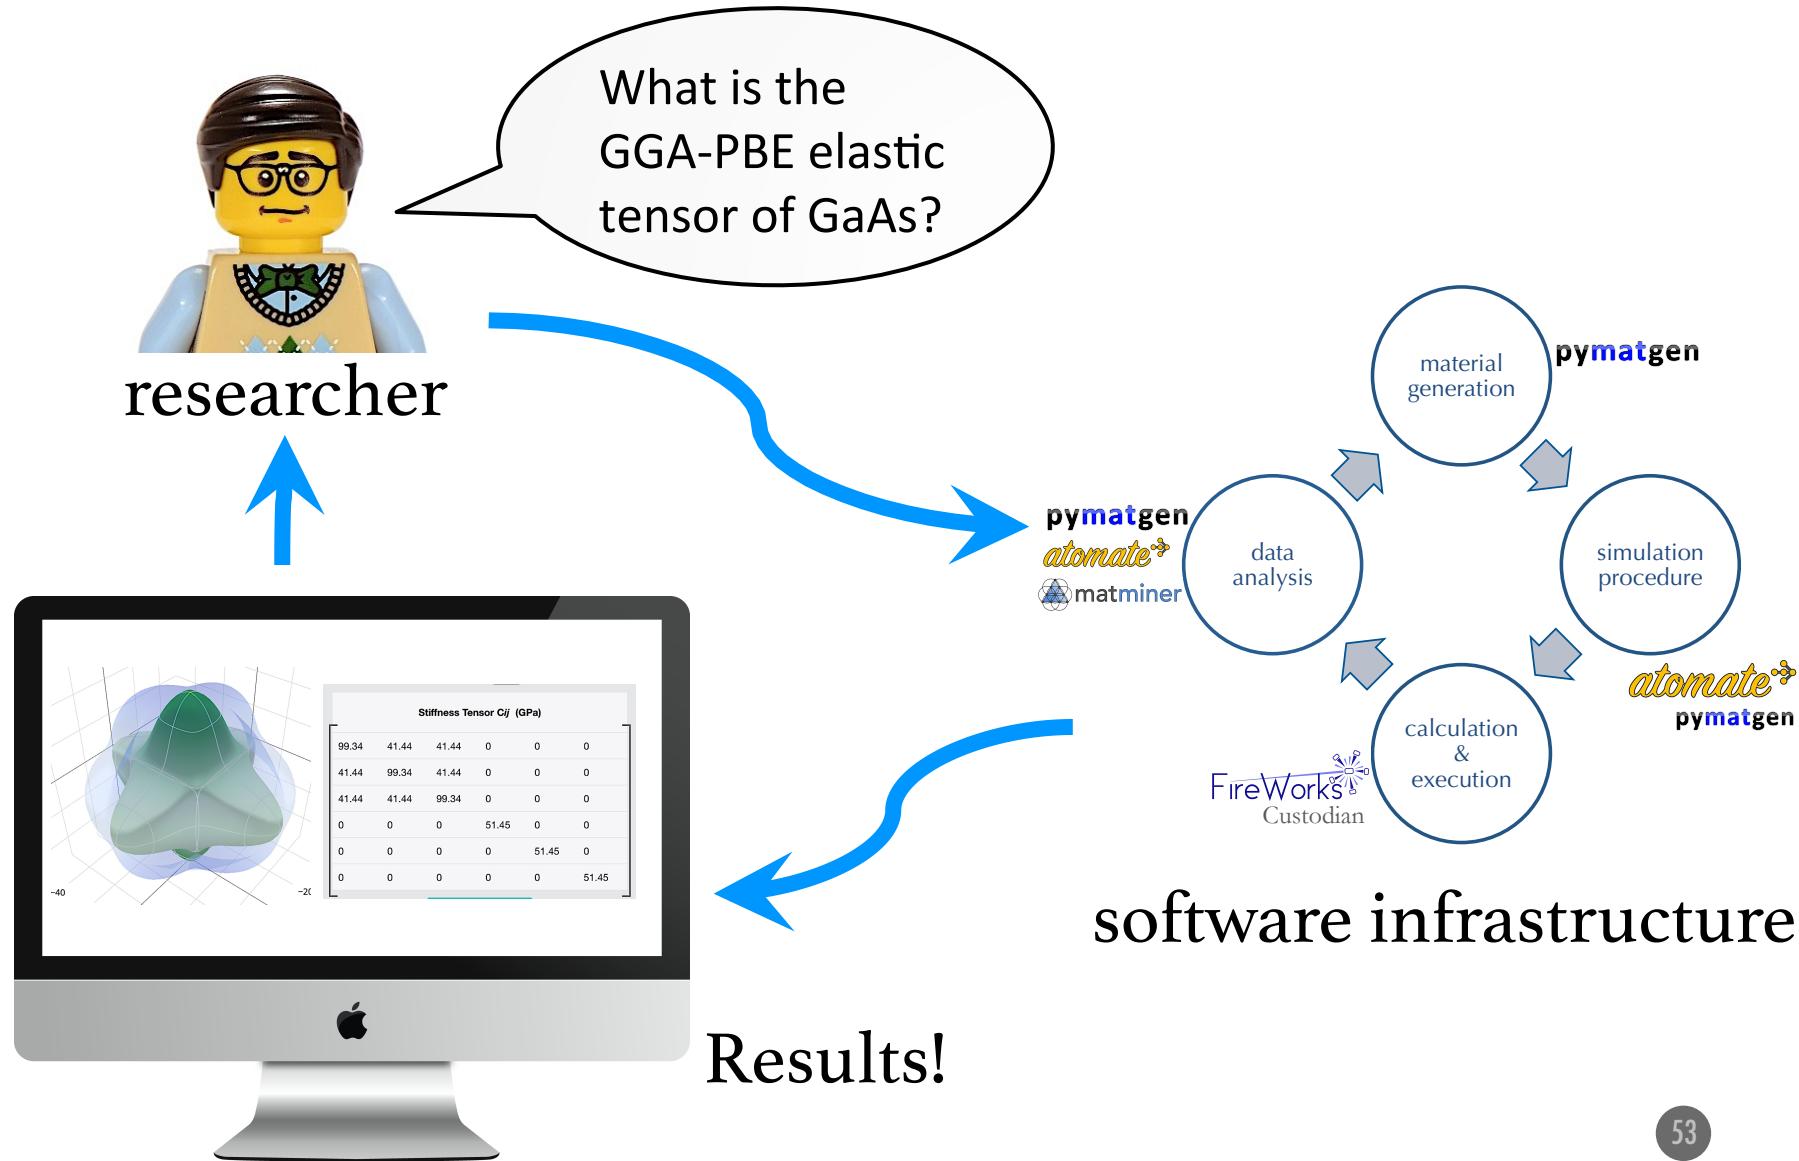
matminer helps enable data mining studies – currently in (open) pre-release



Data analysis: the main ideas

- Atomate generates databases that summarize your calculation results. Additionally The *builders* framework generates document collections that summarize data on a particular material from multiple calculations or from additional information.
- Pymatgen is a powerful software for doing conventional scientific analyses with the computational data
- Matminer, currently in pre-release, is software for doing data mining style analyses of materials

Doing simulations programmatically can become second nature – and very fast / automatic / scalable



Next steps

- This purpose of this presentation was to explain the benefits of the software tools available to you
- You will need to take more steps to actually gain use out of the software
- There are multiple resources to help you with this!

For general information / overview / citing

- There are papers on the software tools for general information (and for citation purposes)

Ong, S. P. et al. Python Materials Genomics (pymatgen): A robust, open-source python library for materials analysis. *Comput. Mater. Sci.* 68, 314–319 (2013).

CONVERGENCE AND COMPUTATION: PRACTICE AND EXPERIENCE
Concurrency Comput. Pract. Exper. (2013)
Published online in Wiley Online Library (wileyonlinelibrary.com) DOI: 10.1002/cpex.3105

FireWorks: a dynamic workflow system designed for high-throughput applications

Amishva Jain^{1,*}, Shye Ping Ong², Wei Chen³, Bharat Mehta³, Xianhui Qu⁴, Michael Kocher⁵, Miftah Baffour⁵, Guido Petruzzelli⁵, Gian-Marco Rignani⁶, Geoffrey Hause⁷, Daniel Gutz⁷ and Kristin A. Persson⁸

¹Environmental Energy and Technology Division, Lawrence Berkeley National Laboratory, Berkeley, CA 94720, USA

²Department of Materials Science and Engineering, University of California San Diego, La Jolla, CA 92093, USA

³Computer Science Department, University of California San Diego, La Jolla, CA 92093, USA

⁴Institute of Condensed Matter and Nanoscience (IMCN), European Theoretical Spectroscopy Facility (ETSF)

Université Catholique de Louvain, Louvain-la-Neuve, Belgium

SUMMARY

This paper introduces FireWorks, a workflow software for running high-throughput calculation workflows at supercomputing centers. FireWorks has been used to complete over 20 million CPU hours of computational chemistry and materials science calculations at the National Research Super-Computing Center. It has been designed to support (i) distributed execution of parallel tasks, (ii) fine-grained performance support for (i) concurrent execution of multiple tasks, (iii) failure detection and correction, (iv) performance and reporting for long-running projects, (v) automated deployment updates, and (vi) dynamic workflow. Our system is built on top of the Apache Mesos framework and the Apache Flink distributed processing system, enabling modern data-driven and high-throughput science applications, and we discuss our implementation strategy on Python and Node.js. (Manuscript accepted 20 March 2013; first published online 27 September 2014)

Received 27 September 2014; Revised 3 March 2015; Accepted 14 April 2015

KEY WORDS: scientific workflow; high-throughput computing; fault-tolerant computing

1. INTRODUCTION

Scientists are embarking on a new paradigm for discovery employing massive computing to generate and analyze large data sets [1]. This approach is spurred by improvements in computer power, faster data acquisition, and advances in sensor technology that have caused scientific computing applications to grow in size and complexity. Some application have leveraged these new capabilities to increase system size, resolution, or accuracy via massive parallelism (e.g., in the simulation of combustion or climate). However, in other areas, researchers employ large amounts of computational power to increase the number of independent calculations, thereby increasing the number of computational cost per calculation. We refer to this mode as queuing (high concurrency of many smaller jobs), and particularly the case where jobs are continually added over long time frames [2].

In this paper, we introduce FireWorks, a workflow system designed for high-throughput applications, such as the Materials Project [3], which can require large amounts of total computing time. For example,

the Materials Project (an effort to generate materials data through simulation) currently uses tens of

*Correspondence to: Amishva Jain, Environmental Energy and Technologies Division, Lawrence Berkeley National Laboratory, Berkeley, CA 94720, USA.
Email: amishva.ljain@lbl.gov

Copyright © 2015 John Wiley & Sons, Ltd.

Jain, A. et al. FireWorks: a dynamic workflow system designed for high-throughput applications. *Concurrent Comput. Pract. Exp.* 22, 5037–5059 (2015).

Mathew, K. et al. Atomate: A high-level interface to generate, execute, and analyze computational materials science workflows. Comput. Mater. Sci. 139, 140–152 (2017).

For installing / usage / examples

- The online documentation is best for practical usage
 - www.pymatgen.org
 - <https://materialsproject.github.io/fireworks/>
 - <https://materialsproject.github.io/custodian/>
 - <https://hackingmaterials.github.io/atomate/>
 - <https://hackingmaterials.github.io/matminer/>
- The online documentation includes installation, examples, tutorials, and descriptions of how to use the code
- If you want to do “everything”, suggest starting with atomate and going from there

For help / questions

- Did you try your best to use the software but are having trouble?
- The various codebases have official channels for getting help
 - <https://groups.google.com/forum/#!forum/pymatgen>
 - <https://groups.google.com/forum/#!forum/fireworkflows>
 - <https://groups.google.com/forum/#!forum/atomate>
 - <https://groups.google.com/forum/#!forum/matminer>
 - for custodian, use Github “issues”
- Some of these lists have resolved >100 user questions, all essentially on volunteer time!

Questions? Comments?

Thanks to everyone that contributed to this software and all those who helped support / justify its development through citation!